

Testing Faster in Continuous Integration Pipelines. An Industrial Case Study

Diego Joshua Martínez Pineda

Supervisor: Mahsa Varshosaz
Submitted: May 2022

IT UNIVERSITY OF COPENHAGEN

Acknowledgments

SYN+ACK¹

¹A TCP joke, because less people would get a UDP one.

Abstract

In this project, we present the evaluation of test selection and prioritization techniques, with data collected from a Continuous Integration (CI) pipeline of the Microsoft *Business Central* project.

We reduce the test prioritization problem to a ranking problem. This is done by computing properties for the given changes to the codebase and using them as input vectors to different ranking algorithms. In contrast to existing research, we use different properties; most notably, the usage of test coverage information.

Having an effective technique for test prioritization in a project is relevant in Continuous Integration pipelines. Their execution can be time-consuming and energy inefficient. Furthermore, these are typically executed multiple times per day in a large team. Speed and efficiency not only benefit the development cycles but overall development costs and energy impact.

Our results show that using coverage information for assigning priorities increases the effectiveness of the prioritization proposed. However, using coverage information as part of the feature vectors reduces its effectiveness and increases the variance of the results for the datasets created.

Based on these results, we argue that with the proposed prioritization algorithms we would be able to reduce the number of tests that were run by 60% using a conservative estimate of the best performing algorithm, which corresponds to an execution time improvement of 90% in average. This suggests that these approaches are well suited for the pipeline under study.

We also give an outline of how this prioritization could be integrated into different stages of the existing pipeline.

Contents

Acknowledgments	iii
Contents	v
1 Introduction	1
2 Background	3
2.1 Business Central	3
2.2 Test Selection and Prioritization techniques	8
3 Method	17
3.1 Outline	17
3.2 Collecting the dataset of CI executions	17
3.3 Representing tests and codebase changes by a vector	18
3.4 Prioritizing Test Executions in a CI job	19
3.5 Training ranking models	21
4 Results and Evaluation	23
4.1 Coordinate Ascent	24
4.2 LambdaMART	24
4.3 MART	26
4.4 Other algorithms	27
5 Conclusion	31
5.1 Threats to validity	31
6 Future Work	33
6.1 Usage of a prioritization technique in the context of the existing system.	33
6.2 Tackling CI optimization from other angles	34
6.3 Further evaluation of test prioritization techniques	34
Bibliography	37
A The AL Language	39
A.1 How does AL look?	39
A.2 AL objects and BC's runtime	40
A.3 AL Tests	43

B Evaluation results	45
----------------------	----

Chapter 1

Introduction

Large software systems are a composition of intertwined units of functionality, often related in complex ways. They can be developed across many years and by large teams. It is no surprise that complexity arises quite easily.

It becomes progressively harder to add new features and improvements while being sure that everything else works as it should. That is, the software remains correct. To tackle this problem of complexity, many strategies and different angles have been studied and implemented by large-scale software systems.

Automated regression testing is one of the main strategies widely adopted by the industry, to ensure that changes to the code will not negatively affect the users of the system.

It consists of developers writing automated tests: code that executes scenarios that a user would experience and asserts that desired properties hold. Afterwards, when changes to the code are done, executing successfully the set of automated tests gives confidence that such scenarios will not be impacted by the change.

In practice, this is an effective technique to keep the quality of the system and allow for its evolution. However, there is a drawback, as the system evolves, more scenarios get added and eventually, it can become time-consuming to run the complete set of automated tests.

Another challenge in software engineering consists in having an effective development process when working with large teams. Many strategies have also been studied, and the industry consensus is to use *Continuous Integration* (CI) processes.

The main idea of the CI software development process is that developers should integrate code changes frequently. When developers integrate their changes, they make other developers on the team aware of their changes. The other developers can then react accordingly if the change conflicts with their work.

Frequently integrating changes to the product increases the risk of making mistakes. For that reason, as a safeguard, the set of automated tests should be run before integrating every change. The process of building and testing the code automatically every time a change is meant to be integrated is called a CI pipeline.

By frequently integrating changes with the regression test suite backing its correctness we ensure quality in every small step.

We can now see that having a time-consuming set of tests to execute is a relevant drawback for teams that want to adopt CI practices. While on one hand, it is desired to

automatically run tests with every small change, a sufficiently large test suite can take a considerable amount of time to execute.

To address this problem, academic research and industry have focused on several angles. Algorithms on Test Selection and Prioritization (TSP) are techniques that aim to automatically propose a relevant sorted subset of tests to execute. These techniques acknowledge that it is not needed to run all the tests, instead, only those which would be more likely to fail first. As the goal of running a test suite is to avoid potential errors to be integrated into the codebase, we can give the developer feedback on errors to be corrected earlier, if found earlier.

For this project, TSP techniques are proposed and evaluated for one of the CI pipelines of the *Business Central* project. *Business Central* is an Enterprise Resource Planning (ERP) software product by Microsoft. The aim is to find an adequate technique for this pipeline and to propose possible directions to take to improve the CI cycle feedback time and effectiveness.

Not only is optimizing CI pipelines relevant to improving the agility of software development, but it also has a positive impact on the energy resources they require. These pipelines are run multiple times a day, and the cost of operating the infrastructure required to keep the system evolving is significant.

In this project, we will follow the same approach as previous research by Bertolino et.al. in [1], transforming this problem into a *ranking* problem. Afterwards, different *Learning to Rank* algorithms will be applied and evaluated. However, in contrast to their work, we use test coverage information for two different parameters of how our ranking datasets are created.

We also outline the different challenges and technical difficulties that adapting academic techniques posed when taking into consideration the practical realities of an industrial environment.

All the code for the different stages of the project can be found in the accompanying repository.¹

¹<https://github.com/mynjj/msc-thesis-public>

Chapter 2

Background

2.1 Business Central

Business Central (BC) is an Enterprise Resource Planning (ERP) software system targeting Small and Mid-sized Businesses (SMB) from Microsoft. Its functionality spans several areas of a company like finance, sales, warehouse management, and many others.

BC allows for Microsoft partners to provide custom extensions that suit customer needs as much as required. This is done through application extensions that developers can write in AL: a Domain Specific Language (DSL) for the application logic. These extensions modify the experience end-customers have with the product, and enhance the product with any custom logic their business may require. Not only partners provide extensions for BC, all the business logic are first-party extensions that Microsoft developers maintain for the core business functionality of the product.

In this thesis project, we will focus on the CI pipeline and regression tests used for the business logic code (also referred to as *application* code). The project has several other pipelines for the different parts of the product, but for the scope of the project, we focus on optimizing the CI pipeline through techniques for test selection and prioritization targeting changes in the AL DSL.

2.1.1 The *application* CI pipeline and the DME system

The BC *application* code follows a traditional CI cycle for development, which we describe for completeness and clarity. Whenever developers complete tasks by making new changes to the *application* code, they create a Pull Request on the Version Control System (VCS) for the desired target branch. Before this request succeeds and the code is merged, certain checks have to be fulfilled, which include the execution of automated tests.

To execute automated tests, a *job* is started on the Distributed Model Executor (DME) system, the internal build system for the product. A job is defined by a set of tasks, defined by scripts that execute the different required stages. Such specification is called the *model* of the job. This set of tasks may have dependencies between each other. We can represent this set of tasks with dependencies as a Directed Acyclic Graph (DAG), where each task corresponds to a node and an edge corresponds to the target depending on the source.

All these tasks and the job definition are defined within the same Version Control repository as code, giving the application developers complete control over the job's ex-




Figure 2.1: Example of tasks that depend on each other for a job execution represented as a DAG.

ecution. A Pull Request triggering the CI pipeline is not the only way a developer can request *jobs*. They can also require them whenever in the development cycle for their convenience.

Some of these tasks correspond to running different groups of tests. Therefore, the cost of executing certain tests belonging to a task depend on the dependencies the task has. We explain by giving an example: Suppose a *job* being executed is defined by a model consisting of the tasks shown in figure 2.1 as the associated DAG for such a set of task-dependencies.

We can see that the task `RunApplicationTests` depends on both `BuildTests` and `BuildApplicationTestsDatabase`. On the other hand, the task `RunSystemTests` only depends on `BuildTests`.

In this example, the DME system could assign two different machines to execute `BuildTests` and `BuildApplicationTestsDatabase`. Let us suppose that `BuildApplicationTestsDatabase` takes longer. When the task `BuildTests` is done, the task `RunSystemTests` can already begin as its only dependency has finished. However, `RunApplicationTests` will have to wait for its other dependency to start execution. In this scenario, running tests in the task `RunApplicationTests` has a higher cost than running tests in `RunSystemTests`.

This means that when executing test selection or prioritization considering the complete set of tests. Some of them entail a higher computing cost depending on the dependencies they have.

The DAG of task-dependencies for the *Application* pipeline of the *Business Central* project is not as simple as the given example. In fact, it is so large that we can not show it meaningfully on a page. However, to satisfy the reader's curiosity, figure 2.2 shows a complete diagram for the metamodel of this pipeline.¹

2.1.2 Test Selection currently on BC *application* tests

In the current CI pipeline of the *application* code of BC, there is a selection method based on test coverage information. We will define more thoroughly the definition of a test selection method in section 2.2, for now, it suffices to think of it as selecting just a subset of tests from the complete set of tests in some meaningful way. In this section, we review how are tests defined for *application* code, and what kind of coverage information we have available when executing tests.

¹A metamodel in this context is a specification for the model that the jobs execute, so it is *smaller*, and easier to show.




Figure 2.2: Associated DAG of the metamodel of job executions of the *Application* pipeline for the *Business Central* product.

```
codeunit 135203 "CF Frcst. Azure AI"
{
    Subtype = Test;
    TestPermissions = NonRestrictive;
    // ...
    [Test]
    procedure AzureAINotEnabledError()
    // ...
    [Test]
    procedure NotEnoughHistoricalData()
    // ...
    [Test]
    procedure FillAzureAIDeltaBiggerThanVariancePercNotInserted()
    // ...
}
```

Figure 2.3: Excerpt of the test codeunit 135203.

Application Tests in AL

We first give an overview of how tests are defined in AL. For a more thorough overview of the AL language and its' role in the *Business Central* system, we refer to Appendix A.

AL organizes the code in *objects*², these *objects* represent different units of functionality for the different features of the product, for example, tables in a database, or pages the user can interact with. A common type of *object* is a *codeunit* which is comprised of different procedures that can be called from any other *AL object*³.

An *application* test is written also in AL, as a *codeunit* with specific annotations for it to be identified by test runner applications. In figure 2.3 you can see an excerpt of a group of test procedures defined under a test codeunit. From a more general perspective, we can think of test codeunits as groups of test scenarios.

Test runners are also implemented in AL. These are programs that run the test codeunits, with different settings, like permission sets to use, or whether or not should the test runner persist changes after test execution, which can be desirable according to the scenarios the developer wishes to test. Currently, the project has two different implementations of test runners, used for different sets of tests, a result of historical legacy. As means of identification, we identify them as *CAL Test Runner* and *AL Test Runner*.

In the current *application* CI pipeline, an *Application Test* task on the job model can use either of these two test runners depending on their needs, or historical legacy. In the current model definition, roughly half of the test tasks are using *CAL Test Runner* and the other half the *AL Test Runner*.

Aside from historic differences, for our purposes, we highlight a significant difference: The *CAL Test Runner* can produce coverage information of tests execution, and the *AL Test Runner* currently used does not.

Coverage information

AL has implemented a language primitive that can be used to record information about which lines of code of which objects were executed. In a broad sense, this is how the

²Note that objects in this context do not correspond to objects in the traditional sense of Object-Oriented Programming.

³*Modules* are a similar concept analogous to *codeunits* in other languages like Python, NodeJS, Rust, Haskell, among others.

```

"Table","4","17","1"
"Table","4","18","1"
...
"Table","14","112","1"
"Table","14","127","1"
...
"Codeunit","28","111","1"
"Codeunit","28","112","1"
"Codeunit","28","115","1"
"Codeunit","28","116","1"
"Codeunit","28","119","1"
"Codeunit","28","122","1"
...

```

Figure 2.4: Excerpt of coverage file for the test codeunit 135203. The columns correspond to: Object Type, Object Id, Line Number and Number of Hits

```

1 codeunit 28 "Error Message Management"
...
106 local procedure GetContextRecID(ContextVariant: Variant; ...
107 var
108     RecRef: RecordRef;
109     TableNo: Integer;
110 begin
111     Clear(ContextRecID);
112     case true of
113         ContextVariant.IsRecord:
114             begin
115                 RecRef.GetTable(ContextVariant);
116                 ContextRecID := RecRef.RecordId;
117             end;

```

Figure 2.5: Some of the lines covered by test codeunit 135203.

test coverage information from the *CAL Test Runner* is produced. After setting up the context for each test, this primitive is used to record activity on each AL object throughout execution.

The *CAL Test Runner* currently has 3 different kinds of output relating to coverage information. We will use information that will allow us to: given a test codeunit executed by the test runner, to have a list of the different *AL objects* and lines within these objects that were run by the test execution.

For concreteness and further illustrate the information available we show an example. In figure 2.3 we have an excerpt of the test codeunit 135203, containing several test procedures. In figure 2.4 we can see a corresponding excerpt of its coverage file. Such coverage was collected for a given state of the codebase.

Each row on the coverage file corresponds to a line being executed when executing test codeunit 135203. The first 3 columns identify such line by giving the Object Type and Object Id the line belongs to, and the Line Number to be considered. The last column corresponds to the number of times the line was executed.

We can see in this example that Codeunit 28 had among others lines 111, 112, 115, and 116 executed. In figure 2.5 we can see such lines.

Test Selection with coverage information

Currently, the CI pipeline performs test selection for tasks executed with the *CAL Test Runner*, for which coverage information of each executed test codeunit is recorded.

The selection technique is based on obtaining the lines and files which the developer is affecting when integrating its change. And finding such lines on the recorded coverage information.

For all those matching lines, the corresponding test codeunit is selected for execution. This effectively selects every test which traversed during its execution the lines being changed by the developer.

However, for tests being run by the *AL Test Runner*, this selection is not available as the corresponding coverage information is not available. We could classify this selection technique as a partial selection based on coverage traces.

As we will explain in section 3.2.1, we will use this coverage information, and add features to the *AL Test Runner* to allow for collecting coverage traces as well.

2.2 Test Selection and Prioritization techniques

Test Selection and Prioritization techniques have a large body of knowledge, the result of extensive focus from both academy and industry.

In this context, the problem of test selection is defined as: given a change to the codebase, and a complete set of tests. Obtain a subset of tests, such that the capacity of fault detection of the test suite is not lost, i. e. if for such change, the test suite execution results in failure, this subset of tests should fail as well. A stronger condition is given by a *safe* selection algorithm, which requires that every failing test case to be included in the selection.

The problem of test prioritization has the same inputs, and it aims to find a sorting for the tests to be executed that prioritizes running the tests that are more likely to fail first.

Intuitively our aim in this problem is that given a change in the codebase, and a complete test suite to execute. Determine a subset (selection) that does not miss any fault-revealing test case, and a sorting (prioritization) that increases the probability of failing first.

It is worth emphasizing that given a prioritization technique, we can induce selection techniques by selecting the first prioritized tests. The size of such selection could be determined by a given size, duration, or other criteria. This is how we obtain the selections evaluated in chapter 4.

2.2.1 Related work

Yoo and Harman in their survey [18] present a detailed overview of techniques in the problems of regression test selection, prioritization, and minimization⁴. In this survey, several foundational works on this area are presented, formal definitions, and metrics to evaluate this problem.

⁴Minimization is a problem not dealt with in this project, it consists of removing superfluous tests from a test suite

It is also worth mentioning the similar techniques found in the literature to the already existing selection technique in *Business Central*. As explained in section 2.1.2 the CI pipeline of interest of the product collects coverage information for some subset of test tasks. Using the information of which files were changed by the developer and the coverage information for each test *codeunit*, a selection is proposed.

This intuitive approach aims to select *modification-traversing* tests [14], this was one of the first approaches studied by Rothermel and Harrold, and widely studied by many others. The approaches differ in how a test is determined to be *modification-traversing*, as some approaches use Control Flow Graphs [13], others use execution traces (Vokolos and Frankl in [16]) and some others use coverage information. Like the work of Beszédes et. al. [2], where they describe populating a coverage database for the C++ components of the WebKit project and identify the changed modules from a given revision.

As a starting point for this project, we reviewed the survey by Pan, et. al. [12] as such survey focused specifically on Selection and Prioritization and techniques using Machine Learning (ML). We initially aimed to find approaches using Reinforcement Learning, as online learning could be beneficial for this use case. In this area, we find the case of Speiker et. al. in [15], where a reinforcement learning agent is proposed using only history of failure and duration of previous iterations.

This leads us to a publication that was the main inspiration and outlined our approach: the proposal by Bertolino et. al. in [1], where they do a more thorough comparison of the effectiveness of Reinforcement Learning approaches, and approaches using ranking algorithms.

Another similar approach is given by Busjaeger and Xie in [4], where they evaluate applying a ranking algorithm for prioritization in the case of Salesforce. We highlight that for the features representing the changes a developer made, they use coverage information. In particular, they propose a coverage score, to reduce the impact of outdated coverage information, which is a reality in large scale environments such as *Business Central*.

2.2.2 Ranking algorithms

As explained in section 2.2.1, previous research has focused on interpreting the problem of test prioritization as a *ranking* problem. In this section, we give a brief overview of this problem, and the algorithms proposed for it. In particular, the techniques we will focus on for this project and how we will interpret the problem of Test Prioritization as an instance of a ranking problem.

In the context of Information Retrieval, the goal of the ranking problem is to obtain relevant resources from a potentially large collection of them for a given information need (query). Ranking algorithms are relevant to different problems. Examples of systems using them are search engines, or recommender systems.

The ranking problem has been extensively studied as it is fundamental for dealing with the information overload that working with computer systems creates.

An approach that has been the subject of extensive research in recent years is *Learning to rank*, a set of ML techniques whose goal is to obtain a *ranking model* out of training data. This model is reused when new queries to the system are given, with the goal of ranking new unseen queries in a similar way as the training data.

A *ranking model* is a mathematical model that given D a collection of documents and a query q , it returns an ordered set of every element of D . Such ordered set is sorted according to some relevance criterion.

In the case of CI cycle optimization with Test Prioritization, we interpret the query q as the change the developer wants to commit to the target branch. The set of documents D corresponds to the complete test suite. Our relevance criterion corresponds to sorting the failing tests first (if any) and the rest of the tests can be sorted through different criteria like duration or test coverage as we will explain in section 3.4.

With this interpretation of the Test Prioritization problem, there is still freedom in two aspects: the representation of the query q for a given codebase change, and the relevance criterion to use. We will vary these parameters as part of our experiments as explained in chapter 3.

We will first describe the different metrics that are used in the ranking literature, and then give a brief overview of the different ranking algorithms used for completeness.

Metrics for the ranking problem

Research on ranking algorithms has given a diverse set of metrics to compare and evaluate rankings proposed by these algorithms. Note that these metrics in our context are used for training and not for evaluating the proposed rankings. This is because it is more meaningful to evaluate with metrics specific to the desired features of the Test Prioritization problem. We will further expand on metrics used for evaluation in section 2.2.3.

A common metric used to evaluate ranking algorithms is the Discounted Cumulative Gain (DCG). $DCG@k$ is defined for k the truncation level as:

$$DCG@k = \sum_{i=1}^k \frac{2^{l_i} - 1}{\log(i + 1)}$$

Where l_i is the relevance given to the i -th result. As we see, this metric increases when the first values are given a high relevance as expected. In contrast, high priority values encountered later are penalized by $\log(i + 1)$.

The truncation level just limits the considered documents for this metric.

$NDCG@k$ is the Normalized version of $DCG@k$, to do so it compares against the ideal ranking for that query and computes its corresponding $DCG@k$, let us call it $IDCG@k$:

$$NDCG@k = \frac{DCG@k}{IDCG@k}$$

The Mean Average Precision (MAP) is based on binary classification metrics. Traditionally precision and recall are widely used for binary classification. In the context of Information Retrieval, *precision* refers to how many documents marked as relevant are relevant by our prediction, *recall* refers to how many of the relevant documents were retrieved from all the relevant documents in the query.

The average precision (AveP) represents the area under the curve of a precision-recall plot when considering the first ranked elements:

$$AveP = \sum_{k=1}^n P(k) |R(k) - R(k - 1)|$$

Where $P(k), R(k)$ are the precision and recall obtained for the first k results.

While our proposed prioritizations to label the dataset are not binary (see section 3.4), they can be considered binary by giving a cutoff point for the assigned relevance.

The Expected Reciprocal Rank (**ERR@k**) metric, was proposed in [5] by Chapelle, et. al.. It is designed to take into consideration the relative ordering between ranked results. In contrast to **DCG**, it does not give the same gain and discount to a fixed position. It is defined by:

$$\sum_{r=1}^k \frac{1}{r} R_r \prod_{i=1}^{r-1} (1 - R_i)$$

Where, for l_m the maximum priority value of the ranking:

$$R_i = \frac{2^{l_i} - 1}{2^{l_m}}$$

While this is the least intuitive of the metrics, it correlates better with search engine applications. It is based on modeling the probability of a user finding its query at a given document position.

As part of our experiments described in chapter 3, we vary the metrics used to train the ranking algorithms.

We will now briefly describe the different ranking algorithms explored.

Coordinate Ascent

Coordinate Ascent is a general *optimization* method, it is based on iterations defined by maximizing the given function f when fixing all coordinates but one. Formally, the k -th iteration, has as i -th component:

$$x_i^{k+1} = \arg \max_{t \in \mathbb{R}} f(x_1^k, \dots, x_{i-1}^k, t, x_{i+1}^k, \dots, x_n^k)$$

This method has similar convergence conditions as gradient descent. It was first proposed as a ranking method by Metzler and Croft in [10], and it has successfully been applied to different ranking problems. We give a short overview of how this optimization method is used as a ranking method, but we refer to the above-mentioned article for a detailed explanation.

Among other things that differ from the traditional optimization, in [10] they propose other constraints to the scoring functions and transformations to reduce the parameter space being optimized to find values on a simplex.

To make it more concrete, the ranking is induced by a scoring function S , for a document D and query Q of the following form:

$$S(D; Q) = \Omega \cdot f(D, Q) + Z$$

For free parameters to optimize Ω , the feature vector $f(D, Q)$ and a constant Z . The free parameters are positive values such that they sum up to 1. In the ranking library used for the implementation in this project, Ranklib, Z is set to zero.

Notice that this scoring function is not the function being optimized, but the ranking evaluation metric.

MART

Regression algorithms using gradient boosting were first proposed by Friedman in [8], Multiple Additive Regression Trees (MART) is a gradient boosting technique further developed by Friedman and Meulman in [9].

Again, we give a very shallow explanation, as a general introduction. We refer the reader to the articles cited above and Machine Learning literature for a more complete explanation.

In general, this is a regression technique that approximates a function by minimizing some related loss function. The idea is to use a linear combination of M different models h_i (called weak learners):

$$F(x) = \sum_i^M \gamma_i h_i(x)$$

The idea is to fit a regression tree to approximate the target function and use the next regression tree to approximate the residuals of the first. Afterwards, greedily compute a scale for the weak learner that minimizes the loss function.

These residuals approximate the gradient of the loss function, effectively making this method follow the same rationale as gradient descent to minimize the loss function.

For *pointwise* ranking algorithms, regression algorithms such as MART can be used to rank by regressing a relevance function for each of the documents to rank, minimizing some of the different training metrics presented in section 2.2.2.

LambdaMART

LambdaMART takes its name from its constituents: LambdaRank and MART. In the previous section, we explained how MART works for general regression and ranking.

In contrast to *pointwise* algorithms that used a relevance function to produce a ranking, *pairwise* algorithms like LambdaMART aim to produce a comparison function between documents.

For the case of this family of algorithms, the aim is to obtain a function that given two documents x_i, x_j , obtains the probability that for a given query q , x_i is ranked with higher than x_j : P_{ij} . With such a comparison function, sorting of the complete set of features can be performed.

To do so, the trained model is a function f that only takes as input a feature x_i , and outputs a real value $f(x_i)$. To obtain the probability of the pairwise comparison a logistic function is used:

$$P_{ij} = \frac{1}{1 + e^{-\sigma(f(x_i) - f(x_j))}}$$

The loss function used is the cross-entropy measure. Minimizing through gradient descent is the idea behind predecessors of this algorithm like RankNet and LambdaRank.

In LambdaMART this gradient is not computed but predicted by boosted regression trees. We refer to [3] for a more thorough explanation of how these gradients are reduced to scalars subject to learning models.

The method has been successfully applied in diverse ranking applications, and in particular, it performed the best in the analysis given by Bertolino, et. al. in [1] on Test Prioritization.

RankBoost

First proposed by Freund, Yoav, et. al. in [7]. As MART, it uses *boosting*, i. e. combining several *weak* learners into a single model.

Each of these learners predicts a relevance function, and therefore a ranking. For training, a distribution D over $X \times X$ is required. Where X is the documents on a query. For this reason, this method is $O(|X|^2)$ in memory, which can restrict the applicability of this algorithm.

Each learner updates the distribution D , emphasizing the pairs that are more relevant for the algorithm to properly order. The final relevance function then becomes a linear combination of each of these learners. We refer to the cited article for more details and further reading.

AdaRank

First proposed in [17] by Xu and Li, it was designed to directly minimize Information Retrieval (IR) performance measures like MAP and NDCG. It is based on AdaBoost, a binary classifier also based on *boosting*, obtaining a model from the linear combination of several *weak* learners.

On each iteration, it maintains a distribution of weights assigned to each of the training queries. Such distribution is updated, increasing the values of the queries that the weak learner ranks the worst. In this way, subsequent learners can focus on those queries for the next rounds.

The weak learners they propose are linear combinations of the metrics to minimize and the weight distribution.

2.2.3 Metrics for evaluating Test Selection and Prioritization

While the problem of ranking has been widely studied and metrics have been proposed for evaluating it. It is more meaningful for our purposes to evaluate the resulting prioritization with metrics in the context of regression testing.

In this section we expand upon some of the metrics previously proposed in the literature, to evaluate the problems of Test Selection and Prioritization.

Test selection execution time

For evaluating selection algorithms, a natural approach is to measure the time it takes to run the subset. To make this metric test suite independent a ratio is used:

$$t_x = \frac{t_S}{t_C}$$

Where t_S is the time taken to execute the selection and t_C is the time taken to execute the complete test suite.

Inclusiveness

In test selection, we do not only rely on the execution time for evaluation. Consider an arbitrarily small subset selected, it would yield good results, but potentially it could also miss some fault-revealing test cases.

To consider this, inclusiveness is introduced:

$$i = \frac{|S_F|}{|T_F|}$$

Where S_F is the set of fault-revealing test cases from a selection, and T_F is the set of test faults in the complete test suite. For completeness, we define $i = 1$ when there are no test faults in the given change.

A *safe* test selection algorithm [13] always has $i = 1$. As every fault-revealing test is included in the selection.

Selection size

On the other hand, high inclusiveness could also be a sign of over-selecting test cases. For example, selecting the whole test suite trivially has $i = 1$. To have a measure of how big the selection is, *selection size* was proposed:

$$ss = \frac{|S|}{|T|}$$

A good selection algorithm strives for having a small selection size, while high inclusiveness.

Time to the first failure

Likewise, another time-related metric of interest for prioritization is the time it takes to reach the first failure:

$$t_{ff} = \frac{t_F}{t_C}$$

Where t_F is the time taken to reach the first failure for the proposed prioritization.

Normalized Average of the Percentage of Faults Detected

For prioritization, only focusing on time to get to the first failure is skewed. As one could have detected the first failure soon while prioritizing the rest of the failing test cases with low priority.

To overcome this, a widely used metric, first proposed by Elbaum, Malishevsky, and Rothermel in [6] is the Average of the Percentage of Faults Detected (APFD).

Normalizing this metric to the number of failures detected allows considering cases where no selected test case was failing. This is useful for evaluating prioritizations that had previously some selection criteria applied.

It is defined by:

$$NAPFD = p - \frac{\sum_{i=1}^m TF_i}{nm} + \frac{p}{2n}$$

For p the ratio of detected faults on a selection against the total amount of faults, n the number of test cases, m the total number of faults, and TF_i the ranking of the test case that revealed the i -th failure.

It represents the proportion of the test failures detected against each executed test. We aim for this value to be close to 1, representing that the accumulated amount of failures detected is obtained early with the prioritization.

Chapter 3

Method

3.1 Outline

In section 3.2, we describe how we obtained the information from the build system of *Business Central*, along with the required changes to collect coverage information.

In section 3.3, we describe how we obtain features for the collected changes and test runs. These features are used to numerically represent a test run in a CI job, for them to be the input of the ranking algorithms.

Afterwards, in section 3.4, we describe how we label the training dataset. That is, the different criteria used to assign priorities to the given tests.

Finally, in section 3.5, we describe the training of the different ranking models, with different datasets.

3.2 Collecting the dataset of CI executions

As the initial step of our project, we collect information about the CI executions from data stored by the DME build system and the history of the repository.

For each CI job, the information extracted was:

- Execution model with the tasks that the DME system used as input.
- Job execution properties: duration, result, and date.
- For each of the tasks executed by the job, information on properties: duration, result, and date.
- Other meta information to identify the job in the VCS.
- For each of the *application test* tasks, the result of each procedure run for each of the test units, along with information on the duration of its execution.
- Comparison with the last merge from the target branch: path and directories where changes were made, type of changes performed, and the content of modified files.

The aim was to collect enough properties to represent the changes a developer made, along with data to evaluate the prioritized tests.

3.2.1 Coverage information for test runners

The information listed in the previous section was collected from real operation data of the pipeline. As explained in section 2.1.2, there are two different implementations of test runners that the tasks may use.

As explained in section 2.1.2, in the current pipeline, tests run with the *CAL test runner* do have coverage information. However, that is not the case for tests that use the *AL test runner*.

A complete coverage report for all the application tests was not initially available. However, as part of this project, modifications to the *AL test runner* were done to allow for collecting the same kind of information¹. However, these changes were not integrated into the pipeline. Instead, the changes made to the test runner were run against snapshots of the codebase in a given time.

It has been discussed previously in research how coverage information may be outdated and hard to maintain [1]. This is partly true in our case as well, however, we acknowledge that the information given by coverage can be valuable for our problem.

In [4] a more robust approach to using coverage information is proposed, by defining a coverage score. A feature like coverage score mitigates the lack of accuracy of the coverage information. As additional mitigation to this problem we introduce windows to compute such coverage scores as it will be shown in section 3.3.

3.2.2 The collection process

As a general overview, over roughly a week period, real CI jobs in this pipeline were collected. Sporadically between these jobs, custom jobs were executed with the required changes to the *AL test runner* to collect coverage information.

For the results presented in this work, two of these coverage collecting jobs were executed and retrieved, and 172 CI jobs were collected.

Whenever coverage information is required to compute the features of a given CI job, the coverage information used will be the closest earlier collected one.

The scale of the collected information limited the number of jobs we were able to collect.²

3.3 Representing tests and codebase changes by a vector

For each of the collected CI jobs, we obtain features to characterize the changes made by the developer, for each of the executed tests in the job.

In this section, we explain the different kinds of information obtained to characterize them.

3.3.1 File changes

We use the following quantities to represent the changes done for each object:

¹The changes done were based on previous work by Nikola Kukrika (nikolak@microsoft.com)

²The information on CI job executions amounted to 6GB of data and the coverage data to 44GB.

- Amount of new AL tables.
- Amount of new AL objects.
- Amount of modified AL objects.
- Amount of removed AL objects.
- Amount of changed tests.
- Amount of non-AL file changes.
- Amount of added lines to AL objects.
- Amount of removed lines from AL objects.

3.3.2 Test history properties

For each of the tests in a job we add the following historic properties:

- Proportion of times the test has failed within the available data.
- From the past k job executions, the proportion of times the test has failed.
- Average duration of this test within the available data.

3.3.3 Coverage properties

Using the most recent coverage information collected previous to the given job, we compute:

- Ratio between lines covered by that test and the average.
- The number of files changed that were covered by that test.
- The number of lines changed that were covered by that test within different *windows*.

The lines changed were not matched exactly but by *windows*. If a change on a line nearby was covered, it was counted for such properties. We added features for different window sizes.

The aim of such is to reduce the impact of having outdated coverage information.

In section 6.3 we discuss some of the different features that future work could consider found in the literature.

3.4 Prioritizing Test Executions in a CI job

As part of our training dataset for the ranking algorithms, we need to associate a ranking to each of the tests executed in a job. This ranking is what we desire the algorithm to learn. This is the process commonly referred to as *dataset labeling* in the context of ML.

It is worth emphasizing that this prioritization is not the one against which the results will be evaluated, as this would be biased. The evaluation will be given only by the TSP metrics presented in section 2.2.3.

A ranking can be defined by a priority function (also called relevance function): when each test case is assigned a priority, this value can be used to produce a ranking where such priority values are in increasing order.

We outline two different approaches taken to assign priority functions.

3.4.1 Failure and Duration decreasing exponential priority

In [1] they propose a priority function for each test case, based on its duration and outcome. They define a score for the i -th test case R_i by:

$$R_i = F_i + e^{-T_i}$$

Where $F_i = 1$ if the test fails, 0 otherwise, and T_i the duration of its execution.

By design, this score ranks first failing tests and breaks ties via their execution duration. We also notice something else, by this being an exponentially decreasing function, changes in duration have a larger effect for tests with small duration, than test executions with large duration.

3.4.2 Coverage discrete priorities

We propose a discrete priority function using coverage information.

Given the i -th test, we define L_i as the number of lines covered by it.

For a given job executing a subset of tests τ , we define $\mu_{L,\tau}$ to be the mean of all values $\{L_i\}_{i \in \tau}$. For such job, our proposed priority function, prioritizes in the following order:

- Failing Tests
- New or modified tests in the job
- Tests that have no coverage information
- Tests covering changed lines, where $L_i \geq \mu_{L,\tau}$
- Tests covering changed lines, where $L_i < \mu_{L,\tau}$
- Tests where $L_i \geq \mu_{L,\tau}$
- Tests where $L_i < \mu_{L,\tau}$

It is worth reminding the reader, that from such a job we do not have the exact coverage information but the most recent, previously collected. As we described in section 3.2.

We created training datasets with both of these priority functions.

3.5 Training ranking models

We performed training of the different described ranking algorithms for each of the training datasets created. For implementation, we used Ranklib 2.17. The training was performed with the High-Performance Cluster from the IT University of Copenhagen.

For each training performed, we varied some of the available hyperparameters and also varied with the training metrics presented in section 2.2.2.

Additionally, for algorithms based on regression trees like MART and LambdaMART, we varied the number of trees for gradient boosting.

We varied the training parameters to get the best performing configuration for each approach as will be presented in chapter 4.

As explained in previous sections, we produced different datasets changing the prioritization criteria and the features that describe each change. As means of identification for evaluation, we name them as described in table 3.1.

Dataset name	Prioritization	Features of each test
EP-NCI	Decreasing exponential as in [1]	Every feature in section 3.3 except coverage related
EP-CI	Decreasing exponential as in [1]	Every feature in section 3.3
CP-NCI	Coverage based as described in section 3.4.2	Every feature in section 3.3 except coverage related
CP-CI	Coverage based as described in section 3.4.2	Every feature in section 3.3

Table 3.1: Naming of the different datasets created.

Each of these datasets consisted of 172 queries corresponding to the collected CI jobs, from which 20% of the failing test jobs were used as the validation dataset. Each of these queries contains around 18000 tests.

It is worth comparing the size of this dataset with the provided datasets from the literature. In [15] the *Paint Control* dataset consists of 180 cycles with an average of 65 tests each. In [1] their dataset consists of 522 cycles with an average of 22 tests each, which they created from the Apache Commons projects commit history.

Chapter 4

Results and Evaluation

First, we use the NAPFD metric as defined in section 2.2.3 to compare the effectiveness of the rankings for the Test Prioritization problem.

In this section, we present for each ranking algorithm used, how the different criteria used influence the NAPFD behavior. To emphasize, the varying criteria for the experiments were:

- Priority function used on datasets.
- Usage of coverage information
- Training metric used for ranking algorithms.
- For MART and LambdaMART, the number of trees used.

For each algorithm, we present box plots of the distribution of this metric for some of these configurations.

Additionally, for each of these ranking algorithms, we induce selection algorithms by taking the first proposed results. We take the first proposed results based on the following criteria:

- An strictly safe selection **S-SEL**: we take the first results such that for all the evaluated jobs, every test failure is included.
- An above 80% average selection **80-SEL**: the first results such that the average of its *inclusiveness* is at least 80%.
- An above 50% average selection **50-SEL**: likewise, we include tests until the average inclusiveness is at least 50%.

We obtain these selections by iterating in increments of 10%.

For a complete set of values shown in the distribution plots and more evaluation metrics, see appendix B and the evaluation data provided in the accompanying repository.

4.1 Coordinate Ascent

Using different training metrics did not have a large impact on the behavior of this algorithm. However, the highest average and lowest variance of the NAPFD of different datasets ranked was obtained with the **NDCG@30** training metric.

Additionally, results show that using coverage information as part of the features characterizing a change results in lower values of NAPFD. However, for the coverage priority function, the NAPFD values were consistently higher.

The best results for this algorithm, across the different metrics, were obtained by datasets that do not use coverage information in the features, but that use the coverage prioritization proposed. In the worst case for these datasets, the NAPFD of the proposed test ranking can be as low as 67%, but as high as 99%.

Figures 4.1 to 4.1 show the box plot of the distribution of NAPFD values for the jobs in the validation dataset for the different training metrics.

Additionally, for **NDCG@30**, we present the distribution of values across each different dataset of the *time to first the failure* metric in figure 4.1.

We observe that only in the case of the dataset with coverage information and coverage prioritization, the proposed prioritization achieves a value as high as 88% of the total execution time with an average of 14%. For the remaining datasets, the first failure is detected at most in the first 2% of the total execution time.

Regarding the induced selections, for the training metric **NDCG@30** and **CP-CI** dataset the following results were obtained:

- **S-SEL**: A safe selection was achieved with a selection size of 40%, corresponding to 10% of the execution time.
- **80-SEL**: A selection with average inclusiveness over 80% was achieved with a selection size of 40%, corresponding to 10% of the execution time.
- **50-SEL**: A selection with average inclusiveness over 50% was achieved with a selection size of 10%, corresponding to 3% of the execution time.

4.2 LambdaMART

For LambdaMART, using the **MAP** metric, consistently ranked with a NAPFD value of 66% for any of the different trained datasets, as it can be seen in figure 4.2. With the **ERR@10** metric, the training did not converge, resulting in an invalid model.

Apart from these two metrics, the other training metrics behaved similarly across the different datasets. The best NAPFD value was obtained with the **DCG@10** metric.

Regarding the number of trees used for gradient boosting, the best performing values were obtained with 20 trees.

For this algorithm, using the coverage prioritization proposed yielded better NAPFD values. When using the exponential prioritization, using coverage information as features increased the NAPFD average and reduced its variance for the majority of the experiments.




Figure 4.1: Distribution across the different datasets of NAPFD values for the Coordinate Ascent algorithm using the DCG@10 training metric.




Figure 4.2: Distribution across the different datasets of NAPFD values for the Coordinate Ascent algorithm using the MAP training metric.




Figure 4.3: Distribution across the different datasets of NAPFD values for the Coordinate Ascent algorithm using the NDCG@10 training metric.




Figure 4.4: Distribution across the different datasets of NAPFD values for the Coordinate Ascent algorithm using the NDCG@20 training metric.

In figure 4.2 we can see the comparison of distributions of *time to the first failure*, for the metric DCG@10 and 20 trees. For this configuration, the CP-NCI dataset, which performed better across configurations, yields an average NAPFD value of 86%.

The resulting induced selections are:

- **S-SEL:** A safe selection was achieved with a selection size of 40%, corresponding to 9% of the execution time.
- **80-SEL:** A selection with average inclusiveness over 80% was achieved with a selection size of 40%, corresponding to 9% of the execution time.
- **50-SEL:** A selection with average inclusiveness over 50% was achieved with a selection size of 10%, corresponding to 2% of the execution time.




Figure 4.5: Distribution across the different datasets of NAPFD values for the Coordinate Ascent algorithm using the NDCG@30 training metric.




Figure 4.6: Distribution across the different datasets of times to first failure for the Coordinate Ascent algorithm using the NDCG@30 training metric.




Figure 4.7: Distribution across the different datasets of NAPFD values for the LambdaMART algorithm using the MAP training metric and 30 trees.

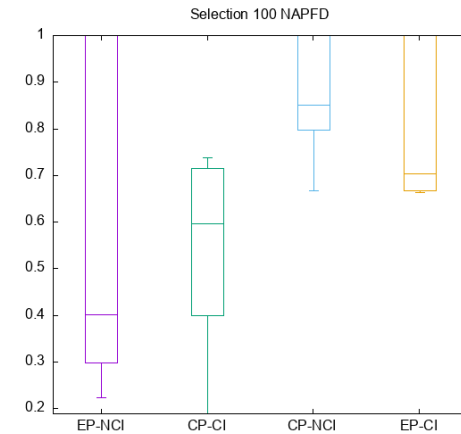


Figure 4.8: Distribution across the different datasets of NAPFD values for the LambdaMART algorithm using the NDCG@10 training metric and 30 trees.

4.3 MART

The results show no significant impact on the training metric used for this algorithm. The best configuration uses 30 trees.

As in the other algorithms, the dataset that provided the best results is the one using coverage information for prioritizing, but not using coverage information on the feature vectors.

For one of the best performing configurations, we can see in figure 4.3 a distribution of the *time to the first failure* values. The best configurations achieved an average NAPFD value of 67%.

For such configuration, the resulting induced selection has similar values as the other algorithms: The resulting induced selections are:

- **S-SEL:** A safe selection was achieved with a selection size of 40%, corresponding to




Figure 4.9: Distribution across the different datasets of NAPFD values for the LambdaMART algorithm using the NDCG@10 training metric and 20 trees.




Figure 4.10: Distribution across the different datasets of NAPFD values for the LambdaMART algorithm using the DCG@10 training metric and 30 trees.




Figure 4.11: Distribution across the different datasets of NAPFD values for the LambdaMART algorithm using the DCG@10 training metric and 20 trees.




Figure 4.12: Distribution across the different datasets of NAPFD values for the LambdaMART algorithm using the NDCG@20 training metric and 30 trees.

9% of the execution time.

- 80-SEL: A selection with average inclusiveness over 80% was achieved with a selection size of 40%, corresponding to 9% of the execution time.
- 50-SEL: A selection with average inclusiveness over 50% was achieved with a selection size of 10%, corresponding to 2% of the execution time.

4.4 Other algorithms

We performed training of AdaRank, with 500 rounds. We varied the different considered metrics as with the other approaches. However, none of the resulting models converged for our dataset.




Figure 4.13: Distribution across the different datasets of NAPFD values for the LambdaMART algorithm using the NDCG@20 training metric and 20 trees.




Figure 4.14: Distribution across the different datasets of NAPFD values for the LambdaMART algorithm using the NDCG@30 training metric and 30 trees.




Figure 4.15: Distribution across the different datasets of NAPFD values for the LambdaMART algorithm using the NDCG@30 training metric and 20 trees.

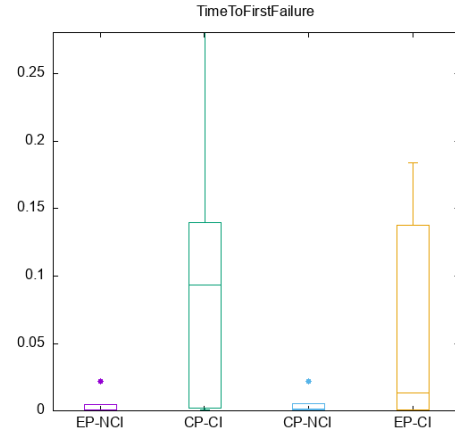


Figure 4.16: Distribution across the different datasets of *time to the first failure* values for the LambdaMART algorithm using the DCG@10 training metric and 20 trees.

Training of RankBoost was also attempted. However, as explained in section 2.2.2, it requires at each stage to keep a distribution D of memory complexity $O(n^2)$ for n the number of tests on each CI job. With the dataset used, such memory complexity exceeded the memory capacity of even the HPC nodes where the training was performed.

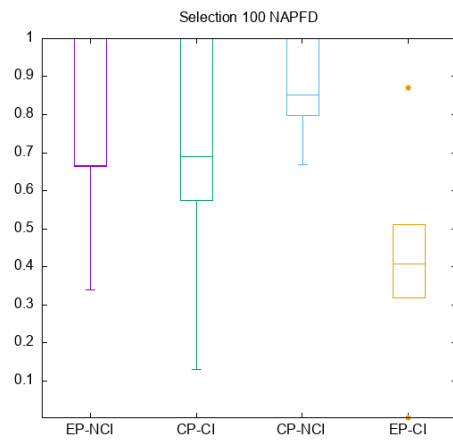


Figure 4.17: Distribution across the different datasets of NAPFD values for the MART algorithm using the DCG@10 training metric and 30 trees.




Figure 4.18: Distribution across the different datasets of NAPFD values for the MART algorithm using the MAP training metric and 30 trees.

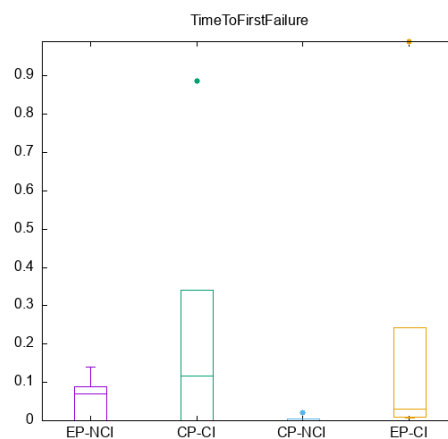


Figure 4.19: Distribution across the different datasets of *time to the first failure* values for the MART algorithm using the DCG@10 training metric and 30 trees.

Chapter 5

Conclusion

In this project, we created different test prioritization datasets with information collected from the CI pipeline of *Business Central*. These datasets included different features to represent changes to a codebase and criteria to prioritize tests.

With these datasets, different ranking algorithms were trained and evaluated. We found that the best performing and most consistent approach was Coordinate Ascent with the training metric NDCG@30 . The induced *safe* selection with this approach reduces the number of tests executed by 60%, and their execution time by 90%.

In contrast to existing approaches to the problem using *Learning to Rank*, we use coverage information. We found that the best-performing dataset used coverage for prioritizing the training dataset.

Existing datasets in research for this problem do not include coverage information. The ranking datasets used for this project can be found in the accompanying repository.

The infrastructure to collect CI history information and related coverage information, as well as the extraction mechanism of features can also be found in the accompanying repository of this project. Such tooling can be used to increase the strength and confidence on the evaluation here presented by adding more samples and extending the time span from which the dataset is collected.

Having a data model for the changes a developer makes in a codebase increases the insight that can be obtained for future analysis of this CI pipeline.

5.1 Threats to validity

Although the results are promising, it is important to consider the threats to the validity of the evaluation. The most relevant threat is a biased dataset.

As explained in section 3.2, the dataset was collected over a week. Also, as previously explained, the size of the collected dataset and test suite restricted the training of the ranking algorithms.

Among some of the reasons this could have biased the evaluation, is that a week is a short time to be representative of the overall test execution dynamics.

For example, a developer working on a feature may re-run the pipeline and continuously fail the same set of related tests. Given that history features are used, the learning algorithms could use this as an indicator of a high priority test. Which may not be true

after the work is done. In an online implementation of these techniques, this problem could be mitigated by periodic retraining of the model.

Chapter 6

Future Work

6.1 Usage of a prioritization technique in the context of the existing system.

This work focused on an offline evaluation of the techniques. For an online implementation, other technical challenges and practical considerations remain.

So far in our discussion, we did not tie how proposing a prioritization with these techniques relates to how they are executed by the DME system explained in section 2.1.1.

In this section, we outline the required changes, and possible strategies to use such rankings.

Recall that the DME system, executes tasks from a given job, by traversing each of the required dependencies defined by the model of the job. A single task may execute multiple test codeunits or test solutions.

Given a ranking proposed by these techniques, we can use the induced selection algorithm to reduce the number of test codeunits that each of these tasks executes.

Furthermore, if the selection results in *job tasks* with no test codeunits to execute, we can remove this task, along with the dependencies that were only required by this task. By doing so recursively we can remove entire paths of the job's execution.

For clarity, see figure 6.1, where a DAG representing the dependencies of a job's execution is given. Filled with black are the test tasks that after the selection algorithm, had no tests to execute. Filled with gray are the tasks that were only required for such tests, and therefore could be removed.




Figure 6.1: Tasks and predecessors removed from the job model.

This would be an effective use of the selection proposed, however, engineering is required to allow for such dynamic changes in a job's execution.

Since test tasks can be performed in parallel by the DME system, a proposed prioritization of all the tests in the job can not be executed as given. Instead, the prioritization can be used locally on each task. For a given test task and the overall prioritization, one can get a local prioritization for the tests belonging to such task. However, the engineering that allows dynamically sorting the test codeunits in a single task is also missing.

Finally, the current build system allows for assigning priorities to the tasks to run, which is taken into consideration when deciding which task to execute next from the set of tasks with completed dependencies. This could be dynamically assigned based on the prioritization of the tests being run.

The main engineering challenge to allow for the usage of these proposals is to make tests *first-class citizens* of the data model proposed by the DME system. Currently, the DME system has no knowledge of the tests being run by a task, as it is instead the responsibility of the task's implementation.

6.2 Tackling CI optimization from other angles

Business Central has several CI pipelines that constitute each of the different development cycles, from different areas that the product has. We studied a single optimization strategy for one of these CI pipelines, namely Test Selection and Prioritization.

However, different strategies exist in the literature, for instance, the usage of test suite minimization to remove superfluous tests from a codebase.

Another approach could be to have learning models to predict the task failure of each of the tasks a job is comprised from. A similar feature vector as obtained for this project representing the change could be part of the training data used to train binary classifiers.

6.3 Further evaluation of test prioritization techniques

Reducing the number of tests considered for each job execution for our training dataset would allow for a larger timespan to be considered in our evaluation. As each job would become more manageable for the learning algorithms and feature extraction infrastructure.

For instance, in this project, we were not able to perform training of RankBoost, as its memory complexity is squared in the number of tests.

This would also increase the confidence and validity of the results presented in this project.

Other approaches that avoid retraining like Reinforcement Learning could be leveraged and be the subject of future work with the produced dataset.

In regards to *Learning to Rank* approaches, more properties could be obtained to give a vector representation of a codebase change and tests.

For instance, in [1], they propose more code metrics that were not in scope for this project. Examples include cyclomatic complexity, number of public or private methods, number of new functions, and others.

Other interesting features were explored in [4], where besides coverage score, they use *text similarity* metrics between the names of the test procedures and paths of the files or their content.

Adding more useful features to this dataset would allow for more robust criteria for the learning algorithms.

Bibliography

- [1] Antonia Bertolino, Antonio Guerriero, Breno Miranda, Roberto Pietrantuono, and Stefano Russo. Learning-to-rank vs ranking-to-learn: Strategies for regression testing in continuous integration. *2020 IEEE/ACM 42nd International Conference on Software Engineering (ICSE)*, pages 1–12, 2020.
- [2] Árpád Beszédes, Tamás Gergely, Lajos Schrettner, Judit Jász, Laszlo Lango, and Tibor Gyimóthy. Code coverage-based regression test selection and prioritization in webkit. *2012 28th IEEE International Conference on Software Maintenance (ICSM)*, pages 46–55, 2012.
- [3] Christopher Burges. From ranknet to lambdarank to lambdamart: An overview. *Learning*, 11, 01 2010.
- [4] Benjamin Busjaeger and Tao Xie. Learning for test prioritization: an industrial case study. *Proceedings of the 2016 24th ACM SIGSOFT International Symposium on Foundations of Software Engineering*, 2016.
- [5] Olivier Chapelle, Donald Metzler, Ya Zhang, and Pierre Grinspan. Expected reciprocal rank for graded relevance. In *Proceedings of the 18th ACM Conference on Information and Knowledge Management, CIKM '09*, page 621–630, New York, NY, USA, 2009. Association for Computing Machinery.
- [6] S. Elbaum, Alexey Malishevsky, and Gregg Rothermel. Test case prioritization: A family of empirical studies. *Software Engineering, IEEE Transactions on*, 28:159–182, 03 2002.
- [7] Yoav Freund, Raj Iyer, Robert E. Schapire, and Yoram Singer. An efficient boosting algorithm for combining preferences. *J. Mach. Learn. Res.*, 4(null):933–969, dec 2003.
- [8] Jerome H. Friedman. Greedy function approximation: A gradient boosting machine. *Annals of Statistics*, 29:1189–1232, 2001.
- [9] Jerome H. Friedman and Jacqueline J. Meulman. Multiple additive regression trees with application in epidemiology. *Statistics in Medicine*, 22, 2003.
- [10] Donald Metzler and W. Bruce Croft. Linear feature-based models for information retrieval. *Information Retrieval*, 10:257–274, 2006.

- [11] Microsoft. Programming in al. <https://docs.microsoft.com/en-us/dynamics365/business-central/dev-itpro/developer/devenv-programming-in-al>. (Accessed on 23/05/2022).
- [12] Rong Pan, Mojtaba Bagherzadeh, Taher Ahmed Ghaleb, and Lionel Claude Briand. Test case selection and prioritization using machine learning: A systematic literature review. *ArXiv*, abs/2106.13891, 2021.
- [13] G. Rothermel and M.J. Harrold. A safe, efficient algorithm for regression test selection. In *1993 Conference on Software Maintenance*, pages 358–367, 1993.
- [14] G. Rothermel and M.J. Harrold. Analyzing regression test selection techniques. *IEEE Transactions on Software Engineering*, 22(8):529–551, 1996.
- [15] Helge Spieker, Arnaud Gotlieb, Dusica Marijan, and Morten Mossige. Reinforcement learning for automatic test case prioritization and selection in continuous integration. *CoRR*, abs/1811.04122, 2018.
- [16] Filippos I. Vokolos and Phyllis G. Frankl. Pythia: a regression test selection tool based on textual differencing. 1997.
- [17] Jun Xu and Hang Li. Adarank: a boosting algorithm for information retrieval. pages 391–398, 07 2007.
- [18] Shin Yoo and Mark Harman. Regression testing minimization, selection and prioritization: a survey. *Softw. Test. Verification Reliab.*, 22:67–120, 2012.

Appendix A

The AL Language

One of the main design aspects of *Business Central* (BC) is that all business logic is written in a custom Application Language (AL), hiding all implementation details dealing with the technology.

Extensions can be provided for the *Business Central* runtime to add business functionality and to fit customers' requirements as much as desired. Such extensions are not only written by Microsoft, as users can customize their experience via extensions as required.

This effectively creates an ecosystem of AL software, for diverse use cases of BC as an ERP.

Adding a language abstraction forces abstraction layers to what is possible within the language and effectively separates business logic from technology implementation specifics. This has allowed the product to evolve throughout the years and experience several changes that with a different design would be harder to achieve, for example: changing the database engine, or migrating to a cloud environment. Additionally, the language continues to evolve and be actively maintained by the team, to keep up with the requirements a modern language infrastructure requires.

Interestingly, this approach has been successfully used by other products as well. For example, another ERP: SAP, has the language ABAP. A custom Domain Specific Language (DSL) resembling COBOL to extend the functionality of the system.

In this section, we give a brief tour of AL, *Business Central*'s custom DSL for modifying its runtime. For a complete reference, we refer to Microsoft's documentation [11].

A.1 How does AL look?

In terms of syntax, AL resembles Pascal. It is an imperative, and procedural language. As a quick overview, we show some of the basic building blocks of the language in the following figures.

As you see, variables are *typed*. The **Record** variable type corresponds to records in a table on the underlying database of the system. As we will explain later, a user can also define the tables to use in this language.

```
var
    myInt: Integer;
    isValid: Boolean;

    Amount := Total * myInt;
```

Figure A.1: Variable declaration in AL.

Figure A.2: Assignment and operations in AL.

```
if x = y then begin
    x := a;
    y := b;
end else
    y := b;
```

Figure A.3: Branching in AL.

```
procedure MyProcedure(Arg1: Integer;
    Arg2: Boolean): Integer
begin
    if Arg2 then
        exit(-Arg1)
    exit(Arg1)
end
```

Figure A.4: Declaring a procedure in AL.

```
<ObjectType> <ObjectID> <ObjectName>
{
    // Definition of the object
}
```

Figure A.5: Syntax to define an AL object.

A.2 AL objects and BC's runtime

Objects in AL are not the general objects as understood in the Object-Oriented Programming paradigm. Instead, they correspond to different units of BC's functionality.

Every code element in AL belongs to some object. To see the syntax of how to declare them, see figure A.5. It requires an `ObjectID` a positive integer, an `ObjectName` a string identifier, and an `ObjectType`.

We will show some of these `ObjectTypes`, and their effect on the runtime.

A.2.1 The Page type

Objects of `Page` type, correspond to interactive interfaces the user experience within the product.

In figure A.6 you can see the AL code used to define the page a user can interact with shown in figure A.7

```
page 379 "Bank Acc. Reconciliation"
{
    Caption = 'Bank Acc. Reconciliation';
    PageType = ListPlus;
    PromotedActionCategories = 'New,Process,Report,Bank,Matching,Posting';
    SaveValues = false;
    SourceTable = "Bank Acc. Reconciliation";
    SourceTableView = WHERE("Statement Type" = CONST("Bank Reconciliation"));

    layout
    {
        // ...
    }
}
```

Figure A.6: Example of an AL page.

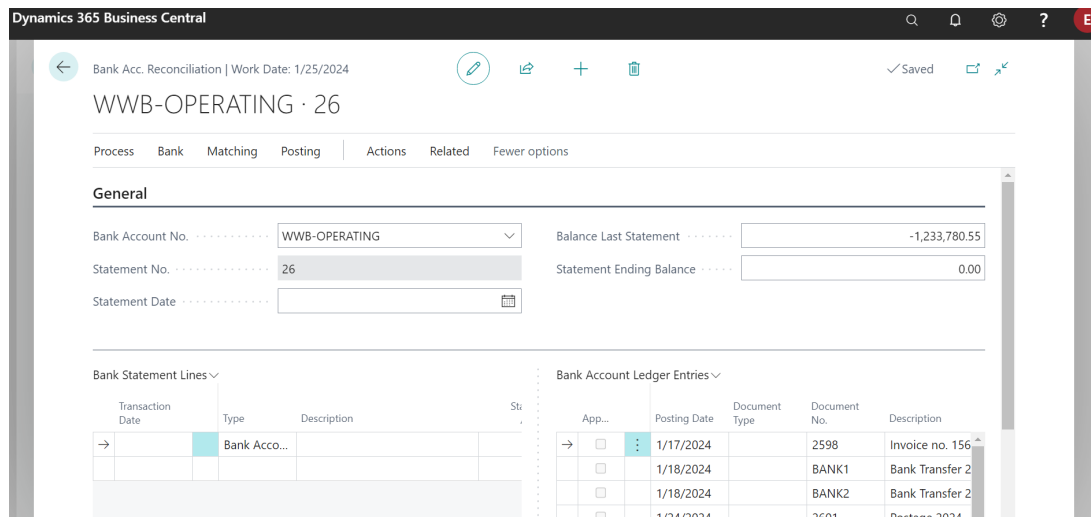


Figure A.7: The interface with which the user interacts, as defined by the AL page from figure A.6.

```
table 273 "Bank Acc. Reconciliation"
{
    Caption = 'Bank Acc. Reconciliation';
    DataCaptionFields = "Bank Account No.", "Statement No.";
    LookupPageID = "Bank Acc. Reconciliation List";
    Permissions = TableData "Bank Account" = rm,
                  TableData "Data Exch." = rimd;

    fields
    {
        field(1; "Bank Account No."; Code[20])
        {
            // ...
        }
        // ...
    }
    // ...
}
```

Figure A.8: Example of an AL Table.

A.2.2 The Table type

Objects of **Table** type, correspond to persistent storage in the system. Defining this object corresponds to creating a table in the underlying SQLServer database.

Having defined a table, a developer can now use **Record** type variables of such table, to manipulate and use this table as required, effectively acting as a data layer abstraction.

In figure A.8 you can see the definition of a table, and in figure A.9 how data can be manipulated by the usage of a **Record** variable.

A.2.3 The Codeunit type

Objects of **Codeunit** type, correspond to logical units of functionality, much like *modules* in other languages. They are a set of procedures that can be called from anywhere else within the AL codebase with certain access modifiers.

In figure A.10 you can see a procedure defined in a codeunit and in figure A.11 an example of it being used from a different AL object.

```
// ...
CurrPage.Update(false);
if not BankAccReconciliation.IsEmpty() then begin
    BankAccReconciliation.Validate("Statement Ending Balance", 0.0);
    BankAccReconciliation.Modify();
end;
// ...
```

Figure A.9: Example of using a record BankAccReconciliation of the type defined by the table in figure A.8.

```
codeunit 18 "Financial Report Mgt."
{
    TableNo = "Financial Report";

    var
        AccSchedLine: Record "Acc. Schedule Line";
        FinRepPrefixTxt: Label 'FIN.REP.', MaxLength = 10, // ...
        TwoPosTxt: Label '%1%2', Locked = true;
    // ...
    procedure XMLExchangeExport(FinancialReport: Record "Financial Report")
    var
        ConfigPackage: Record "Config. Package";
        ConfigXMLExchange: Codeunit "Config. XML Exchange";
    begin
        AddFinancialReportToConfigPackage(FinancialReport.Name, ConfigPackage);
        Commit();
        ConfigXMLExchange.ExportPackage(ConfigPackage);
    end;

    local procedure AddFinancialReportToConfigPackage(FinancialReportName: Code[10]; var ConfigPackage //...
    // ...
}
```

Figure A.10: Example of an AL Codeunit.

```
// ...
trigger OnAction()
var
    FinancialReportMgt: Codeunit "Financial Report Mgt.";
begin
    FinancialReportMgt.XMLExchangeExport(Rec);
end;
// ...
```

Figure A.11: Usage of the procedure defined in figure A.10 in a different AL object


```

codeunit 134141 "ERM Bank Reconciliation"
{
    Permissions = TableData "Bank Account Ledger Entry" = ri,
                  TableData "Bank Account Statement" = rimd;
    Subtype = Test;
    TestPermissions = NonRestrictive;
    // ...
    [Test]
    [Scope('OnPrem')]
    procedure BankAccReconciliationBalanceToReconcile()
    var
        BankAccReconciliation: Record "Bank Acc. Reconciliation";
        GenJournalLine: Record "Gen. Journal Line";
        BankAccReconciliationPage: TestPage "Bank Acc. Reconciliation";
        BalanceToReconcile: Decimal;
        i: Integer;
    begin
        // [SCENARIO 363054] "Balance to Reconcile" does not include amounts from Posted Bank Reconciliations
        Initialize();

        // [GIVEN] Posted Bank Reconciliation A with Amount X
        CreateAndPostGenJournalLine(GenJournalLine, CreateBankAccount);
        CreateSuggestedBankReconc(BankAccReconciliation, GenJournalLine."Bal. Account No.", false);
        LibraryERM.PostBankAccReconciliation(BankAccReconciliation);

        // [GIVEN] Bank Reconciliation B with Amount Y
        for i := 1 to LibraryRandom.RandInt(5) do begin
            CreateAndPostGenJournalLine(GenJournalLine, GenJournalLine."Bal. Account No.");
            BalanceToReconcile += GenJournalLine.Amount;
        end;
        Clear(BankAccReconciliation);
        CreateSuggestedBankReconc(BankAccReconciliation, GenJournalLine."Bal. Account No.", false);

        // [WHEN] Bank Reconciliation B page is opened
        LibraryLowerPermissions.AddAccountReceivables;
        BankAccReconciliationPage.OpenView;
        BankAccReconciliationPage.GotoRecord(BankAccReconciliation);

        // [THEN] "Balance To Reconcile" = Y.
        Assert.AreEqual(
            -BalanceToReconcile,
            BankAccReconciliationPage.ApplyBankLedgerEntries.BalanceToReconcile.AsDEcimal,
            StrSubstNo(
                WrongAmountErr, BankAccReconciliationPage.ApplyBankLedgerEntries.BalanceToReconcile.Caption,
                -BalanceToReconcile));
    end;
    // ...
}

```

Figure A.12: Example of a test codeunit and test procedure.

A.3 AL Tests

Of particular relevance for this project, is how tests are defined in this language. Tests in AL are defined on AL objects of type `Codeunit` with appropriate annotations.

Depending on the scope of a test, these tests can be integration tests or unit tests. See an example of a test in figure A.12.

The test infrastructure required for running these scenarios with different settings is also maintained by the team and written in AL itself. We go to a bit more detail, as required in section 2.1.2.

We have just scratched the surface with this brief introduction, as it is meant to give a general idea of the type of programs and tests that this thesis project centers around.

Appendix B

Evaluation results

For completeness, we present a list of the main metrics obtained for the different algorithms, ranking configurations and hyperparameters considered.

More concretely, we present the distribution metrics of $NAPFD$ and t_{ff} of each different configuration. We also present the size of the induced selections of each configuration and their corresponding execution times.

More values and distribution plots can be found in the accompanying repository of this project.

Table B.1: $NAPFD$ per dataset, algorithm and configuration

Dataset	Algorithm	Training metric	# of trees	$NAPFD$			
				Average	Variance	Minimum	Maximum
CP-CI	coordinateascent	NDCG@10		0.7	0.08	0.17	0.99
EP-CI	coordinateascent	NDCG@10		0.62	0.11	0.31	1
CP-NCI	coordinateascent	NDCG@10		0.86	0.01	0.67	1
EP-NCI	coordinateascent	NDCG@10		0.86	0.01	0.67	1
CP-CI	coordinateascent	DCG@10		0.7	0.07	0.17	0.96
EP-CI	coordinateascent	DCG@10		0.67	0.09	0.12	1
CP-NCI	coordinateascent	DCG@10		0.86	0.01	0.67	1
EP-NCI	coordinateascent	DCG@10		0.86	0.01	0.67	1
CP-CI	coordinateascent	MAP		0.51	0.1	0.17	1
EP-CI	coordinateascent	MAP		0.51	0.1	0.17	1
CP-NCI	coordinateascent	MAP		0.86	0.01	0.67	1
EP-NCI	coordinateascent	MAP		0.86	0.01	0.67	1
CP-CI	coordinateascent	NDCG@20		0.73	0.08	0.17	1
EP-CI	coordinateascent	NDCG@20		0.62	0.11	0.31	1
CP-NCI	coordinateascent	NDCG@20		0.86	0.01	0.67	1
EP-NCI	coordinateascent	NDCG@20		0.86	0.01	0.67	1
CP-CI	coordinateascent	NDCG@30		0.71	0.08	0.17	1
EP-CI	coordinateascent	NDCG@30		0.88	0.01	0.75	1
CP-NCI	coordinateascent	NDCG@30		0.86	0.01	0.67	1
EP-NCI	coordinateascent	NDCG@30		0.86	0.01	0.67	1
CP-CI	lambdamart	NDCG@10	30	0.56	0.04	0.19	0.74
EP-CI	lambdamart	NDCG@10	30	0.79	0.02	0.66	1
CP-NCI	lambdamart	NDCG@10	30	0.86	0.01	0.67	1
EP-NCI	lambdamart	NDCG@10	30	0.56	0.11	0.22	1
CP-CI	lambdamart	NDCG@10	20	0.65	0.01	0.51	0.82
EP-CI	lambdamart	NDCG@10	20	0.68	0.04	0.5	0.99
CP-NCI	lambdamart	NDCG@10	20	0.86	0.01	0.67	1
EP-NCI	lambdamart	NDCG@10	20	0.56	0.11	0.22	1
CP-CI	lambdamart	NDCG@10	10	0.71	0.01	0.61	0.85
EP-CI	lambdamart	NDCG@10	10	0.64	0.04	0.48	0.99
CP-NCI	lambdamart	NDCG@10	10	0.86	0.01	0.67	1
EP-NCI	lambdamart	NDCG@10	10	0.77	0.05	0.33	1
CP-CI	lambdamart	NDCG@10	5	0.71	0.01	0.61	0.85
EP-CI	lambdamart	NDCG@10	5	0.64	0.05	0.48	0.99
CP-NCI	lambdamart	NDCG@10	5	0.86	0.01	0.67	1
EP-NCI	lambdamart	NDCG@10	5	0.56	0.11	0.22	1
CP-CI	lambdamart	DCG@10	30	0.75	0.02	0.51	0.95
EP-CI	lambdamart	DCG@10	30	0.69	0.08	0.38	1
CP-NCI	lambdamart	DCG@10	30	0.86	0.01	0.67	1
EP-NCI	lambdamart	DCG@10	30	0.65	0.08	0.34	1
CP-CI	lambdamart	DCG@10	20	0.71	0.01	0.61	0.85
EP-CI	lambdamart	DCG@10	20	0.69	0.08	0.38	1
CP-NCI	lambdamart	DCG@10	20	0.86	0.01	0.67	1
EP-NCI	lambdamart	DCG@10	20	0.86	0.01	0.67	1
CP-CI	lambdamart	DCG@10	10	0.71	0.01	0.61	0.85
EP-CI	lambdamart	DCG@10	10	0.69	0.08	0.38	1
CP-NCI	lambdamart	DCG@10	10	0.86	0.01	0.67	1
EP-NCI	lambdamart	DCG@10	10	0.56	0.11	0.22	1
CP-CI	lambdamart	DCG@10	5	0.71	0.01	0.61	0.85

Table B.1: *NAPFD* per dataset, algorithm and configuration

EP-CI	lambdamart	DCG@10	5	0.6	0.09	0.33	1
CP-NCI	lambdamart	DCG@10	5	0.79	0.02	0.66	1
EP-NCI	lambdamart	DCG@10	5	0.65	0.08	0.34	1
CP-CI	lambdamart	MAP	30	0.67	0	0.66	0.67
EP-CI	lambdamart	MAP	30	0.67	0	0.66	0.67
CP-NCI	lambdamart	MAP	30	0.67	0	0.66	0.67
EP-NCI	lambdamart	MAP	30	0.67	0	0.66	0.67
CP-CI	lambdamart	MAP	20	0.67	0	0.66	0.67
EP-CI	lambdamart	MAP	20	0.67	0	0.66	0.67
CP-NCI	lambdamart	MAP	20	0.67	0	0.66	0.67
EP-NCI	lambdamart	MAP	20	0.67	0	0.66	0.67
CP-CI	lambdamart	MAP	10	0.67	0	0.66	0.67
EP-CI	lambdamart	MAP	10	0.67	0	0.66	0.67
CP-NCI	lambdamart	MAP	10	0.67	0	0.66	0.67
EP-NCI	lambdamart	MAP	10	0.67	0	0.66	0.67
CP-CI	lambdamart	MAP	5	0.67	0	0.66	0.67
EP-CI	lambdamart	MAP	5	0.67	0	0.66	0.67
CP-NCI	lambdamart	MAP	5	0.67	0	0.66	0.67
EP-NCI	lambdamart	MAP	5	0.67	0	0.66	0.67
CP-CI	lambdamart	NDCG@20	30	0.68	0.07	0.13	0.95
EP-CI	lambdamart	NDCG@20	30	0.64	0.04	0.48	1
CP-NCI	lambdamart	NDCG@20	30	0.86	0.01	0.67	1
EP-NCI	lambdamart	NDCG@20	30	0.56	0.11	0.22	1
CP-CI	lambdamart	NDCG@20	20	0.62	0.02	0.34	0.81
EP-CI	lambdamart	NDCG@20	20	0.64	0.05	0.49	1
CP-NCI	lambdamart	NDCG@20	20	0.86	0.01	0.67	1
EP-NCI	lambdamart	NDCG@20	20	0.56	0.11	0.22	1
CP-CI	lambdamart	NDCG@20	10	0.71	0.01	0.61	0.85
EP-CI	lambdamart	NDCG@20	10	0.66	0.05	0.43	0.99
CP-NCI	lambdamart	NDCG@20	10	0.86	0.01	0.67	1
EP-NCI	lambdamart	NDCG@20	10	0.42	0.1	0	0.99
CP-CI	lambdamart	NDCG@20	5	0.71	0.01	0.61	0.85
EP-CI	lambdamart	NDCG@20	5	0.67	0.04	0.5	0.99
CP-NCI	lambdamart	NDCG@20	5	0.86	0.01	0.67	1
EP-NCI	lambdamart	NDCG@20	5	0.42	0.1	0	0.99
CP-CI	lambdamart	NDCG@30	30	0.69	0.06	0.19	0.94
EP-CI	lambdamart	NDCG@30	30	0.53	0.06	0.18	0.89
CP-NCI	lambdamart	NDCG@30	30	0.86	0.01	0.67	1
EP-NCI	lambdamart	NDCG@30	30	0.56	0.11	0.22	1
CP-CI	lambdamart	NDCG@30	20	0.64	0.01	0.51	0.77
EP-CI	lambdamart	NDCG@30	20	0.6	0.03	0.33	0.93
CP-NCI	lambdamart	NDCG@30	20	0.72	0.11	0.01	1
EP-NCI	lambdamart	NDCG@30	20	0.56	0.11	0.22	1
CP-CI	lambdamart	NDCG@30	10	0.71	0.01	0.61	0.85
EP-CI	lambdamart	NDCG@30	10	0.79	0.02	0.66	1
CP-NCI	lambdamart	NDCG@30	10	0.72	0.11	0.01	1
EP-NCI	lambdamart	NDCG@30	10	0.42	0.1	0	0.99
CP-CI	lambdamart	NDCG@30	5	0.71	0.01	0.61	0.85
EP-CI	lambdamart	NDCG@30	5	0.54	0.07	0	0.86
CP-NCI	lambdamart	NDCG@30	5	0.72	0.11	0.01	1
EP-NCI	lambdamart	NDCG@30	5	0.42	0.1	0	0.99
CP-CI	mart	NDCG@10	30	0.67	0.09	0.13	1

Table B.1: *NAPFD* per dataset, algorithm and configuration

EP-CI	mart	NDCG@10	30	0.42	0.07	0	0.87
CP-NCI	mart	NDCG@10	30	0.86	0.01	0.67	1
EP-NCI	mart	NDCG@10	30	0.71	0.05	0.34	1
CP-CI	mart	NDCG@10	20	0.67	0.09	0.13	1
EP-CI	mart	NDCG@10	20	0.45	0.06	0	0.83
CP-NCI	mart	NDCG@10	20	0.72	0.11	0.01	1
EP-NCI	mart	NDCG@10	20	0.57	0.06	0	0.67
CP-CI	mart	NDCG@10	10	0.68	0.08	0.17	1
EP-CI	mart	NDCG@10	10	0.48	0.07	0	0.86
CP-NCI	mart	NDCG@10	10	0.57	0.06	0.01	0.67
EP-NCI	mart	NDCG@10	10	0.57	0.06	0	0.67
CP-CI	mart	NDCG@10	5	0.68	0.08	0.17	1
EP-CI	mart	NDCG@10	5	0.54	0.07	0	0.88
CP-NCI	mart	NDCG@10	5	0.57	0.06	0.01	0.67
EP-NCI	mart	NDCG@10	5	0.57	0.06	0	0.67
CP-CI	mart	ERR@10	30	0.67	0.09	0.13	1
EP-CI	mart	ERR@10	30	0.42	0.07	0	0.87
CP-NCI	mart	ERR@10	30	0.86	0.01	0.67	1
EP-NCI	mart	ERR@10	30	0.71	0.05	0.34	1
CP-CI	mart	ERR@10	20	0.67	0.09	0.13	1
EP-CI	mart	ERR@10	20	0.45	0.06	0	0.83
CP-NCI	mart	ERR@10	20	0.72	0.11	0.01	1
EP-NCI	mart	ERR@10	20	0.57	0.06	0	0.67
CP-CI	mart	ERR@10	10	0.68	0.08	0.17	1
EP-CI	mart	ERR@10	10	0.48	0.07	0	0.86
CP-NCI	mart	ERR@10	10	0.57	0.06	0.01	0.67
EP-NCI	mart	ERR@10	10	0.57	0.06	0	0.67
CP-CI	mart	ERR@10	5	0.68	0.08	0.17	1
EP-CI	mart	ERR@10	5	0.54	0.07	0	0.88
CP-NCI	mart	ERR@10	5	0.57	0.06	0.01	0.67
EP-NCI	mart	ERR@10	5	0.57	0.06	0	0.67
CP-CI	mart	DCG@10	30	0.67	0.09	0.13	1
EP-CI	mart	DCG@10	30	0.42	0.07	0	0.87
CP-NCI	mart	DCG@10	30	0.86	0.01	0.67	1
EP-NCI	mart	DCG@10	30	0.71	0.05	0.34	1
CP-CI	mart	DCG@10	20	0.67	0.09	0.13	1
EP-CI	mart	DCG@10	20	0.45	0.06	0	0.83
CP-NCI	mart	DCG@10	20	0.72	0.11	0.01	1
EP-NCI	mart	DCG@10	20	0.57	0.06	0	0.67
CP-CI	mart	DCG@10	10	0.68	0.08	0.17	1
EP-CI	mart	DCG@10	10	0.48	0.07	0	0.86
CP-NCI	mart	DCG@10	10	0.57	0.06	0.01	0.67
EP-NCI	mart	DCG@10	10	0.57	0.06	0	0.67
CP-CI	mart	DCG@10	5	0.68	0.08	0.17	1
EP-CI	mart	DCG@10	5	0.54	0.07	0	0.88
CP-NCI	mart	DCG@10	5	0.57	0.06	0.01	0.67
EP-NCI	mart	DCG@10	5	0.57	0.06	0	0.67
CP-CI	mart	MAP	30	0.67	0.09	0.13	1
EP-CI	mart	MAP	30	0.42	0.07	0	0.87
CP-NCI	mart	MAP	30	0.86	0.01	0.67	1
EP-NCI	mart	MAP	30	0.71	0.05	0.34	1
CP-CI	mart	MAP	20	0.67	0.09	0.13	1

Table B.1: *NAPFD* per dataset, algorithm and configuration

EP-CI	mart	MAP	20	0.45	0.06	0	0.83
CP-NCI	mart	MAP	20	0.72	0.11	0.01	1
EP-NCI	mart	MAP	20	0.57	0.06	0	0.67
CP-CI	mart	MAP	10	0.68	0.08	0.17	1
EP-CI	mart	MAP	10	0.48	0.07	0	0.86
CP-NCI	mart	MAP	10	0.57	0.06	0.01	0.67
EP-NCI	mart	MAP	10	0.57	0.06	0	0.67
CP-CI	mart	MAP	5	0.68	0.08	0.17	1
EP-CI	mart	MAP	5	0.54	0.07	0	0.88
CP-NCI	mart	MAP	5	0.57	0.06	0.01	0.67
EP-NCI	mart	MAP	5	0.57	0.06	0	0.67
CP-CI	mart	NDCG@20	30	0.67	0.09	0.13	1
EP-CI	mart	NDCG@20	30	0.42	0.07	0	0.87
CP-NCI	mart	NDCG@20	30	0.86	0.01	0.67	1
EP-NCI	mart	NDCG@20	30	0.71	0.05	0.34	1
CP-CI	mart	NDCG@20	20	0.67	0.09	0.13	1
EP-CI	mart	NDCG@20	20	0.45	0.06	0	0.83
CP-NCI	mart	NDCG@20	20	0.72	0.11	0.01	1
EP-NCI	mart	NDCG@20	20	0.57	0.06	0	0.67
CP-CI	mart	NDCG@20	10	0.68	0.08	0.17	1
EP-CI	mart	NDCG@20	10	0.48	0.07	0	0.86
CP-NCI	mart	NDCG@20	10	0.57	0.06	0.01	0.67
EP-NCI	mart	NDCG@20	10	0.57	0.06	0	0.67
CP-CI	mart	NDCG@20	5	0.68	0.08	0.17	1
EP-CI	mart	NDCG@20	5	0.54	0.07	0	0.88
CP-NCI	mart	NDCG@20	5	0.57	0.06	0.01	0.67
EP-NCI	mart	NDCG@20	5	0.57	0.06	0	0.67
CP-CI	mart	NDCG@30	30	0.67	0.09	0.13	1
EP-CI	mart	NDCG@30	30	0.42	0.07	0	0.87
CP-NCI	mart	NDCG@30	30	0.86	0.01	0.67	1
EP-NCI	mart	NDCG@30	30	0.71	0.05	0.34	1
CP-CI	mart	NDCG@30	20	0.67	0.09	0.13	1
EP-CI	mart	NDCG@30	20	0.45	0.06	0	0.83
CP-NCI	mart	NDCG@30	20	0.72	0.11	0.01	1
EP-NCI	mart	NDCG@30	20	0.57	0.06	0	0.67
CP-CI	mart	NDCG@30	10	0.68	0.08	0.17	1
EP-CI	mart	NDCG@30	10	0.48	0.07	0	0.86
CP-NCI	mart	NDCG@30	10	0.57	0.06	0.01	0.67
EP-NCI	mart	NDCG@30	10	0.57	0.06	0	0.67
CP-CI	mart	NDCG@30	5	0.68	0.08	0.17	1
EP-CI	mart	NDCG@30	5	0.54	0.07	0	0.88
CP-NCI	mart	NDCG@30	5	0.57	0.06	0.01	0.67
EP-NCI	mart	NDCG@30	5	0.57	0.06	0	0.67

Table B.2: t_{ff} per dataset, algorithm and configuration

Dataset	Algorithm	Training metric	# of trees	t_{ff}			
				Average	Variance	Minimum	Maximum
CP-CI	coordinateascent	NDCG@10		0.15	0.11	0	0.89
EP-CI	coordinateascent	NDCG@10		0.02	0	0	0.14
CP-NCI	coordinateascent	NDCG@10		0	0	0	0.02
EP-NCI	coordinateascent	NDCG@10		0	0	0	0.02
CP-CI	coordinateascent	DCG@10		0.17	0.1	0	0.89
EP-CI	coordinateascent	DCG@10		0.13	0.11	0	0.88
CP-NCI	coordinateascent	DCG@10		0	0	0	0.02
EP-NCI	coordinateascent	DCG@10		0	0	0	0.02
CP-CI	coordinateascent	MAP		0.43	0.1	0	0.89
EP-CI	coordinateascent	MAP		0.43	0.1	0	0.89
CP-NCI	coordinateascent	MAP		0	0	0	0.02
EP-NCI	coordinateascent	MAP		0	0	0	0.02
CP-CI	coordinateascent	NDCG@20		0.14	0.11	0	0.89
EP-CI	coordinateascent	NDCG@20		0.02	0	0	0.14
CP-NCI	coordinateascent	NDCG@20		0	0	0	0.02
EP-NCI	coordinateascent	NDCG@20		0	0	0	0.02
CP-CI	coordinateascent	NDCG@30		0.14	0.11	0	0.88
EP-CI	coordinateascent	NDCG@30		0	0	0	0.01
CP-NCI	coordinateascent	NDCG@30		0	0	0	0.02
EP-NCI	coordinateascent	NDCG@30		0.01	0	0	0.02
CP-CI	lambdamart	NDCG@10	30	0.3	0.09	0	0.89
EP-CI	lambdamart	NDCG@10	30	0.04	0	0	0.14
CP-NCI	lambdamart	NDCG@10	30	0	0	0	0.02
EP-NCI	lambdamart	NDCG@10	30	0.06	0	0	0.14
CP-CI	lambdamart	NDCG@10	20	0.23	0.07	0	0.7
EP-CI	lambdamart	NDCG@10	20	0.08	0.01	0	0.2
CP-NCI	lambdamart	NDCG@10	20	0.01	0	0	0.02
EP-NCI	lambdamart	NDCG@10	20	0.06	0	0	0.14
CP-CI	lambdamart	NDCG@10	10	0.1	0.01	0	0.28
EP-CI	lambdamart	NDCG@10	10	0.13	0.01	0.01	0.25
CP-NCI	lambdamart	NDCG@10	10	0	0	0	0.02
EP-NCI	lambdamart	NDCG@10	10	0	0	0	0.02
CP-CI	lambdamart	NDCG@10	5	0.1	0.01	0	0.28
EP-CI	lambdamart	NDCG@10	5	0.12	0.03	0.01	0.45
CP-NCI	lambdamart	NDCG@10	5	0	0	0	0.02
EP-NCI	lambdamart	NDCG@10	5	0.06	0	0	0.14
CP-CI	lambdamart	DCG@10	30	0.17	0.06	0	0.7
EP-CI	lambdamart	DCG@10	30	0.05	0.01	0	0.18
CP-NCI	lambdamart	DCG@10	30	0	0	0	0.02
EP-NCI	lambdamart	DCG@10	30	0.04	0	0	0.14
CP-CI	lambdamart	DCG@10	20	0.1	0.01	0	0.28
EP-CI	lambdamart	DCG@10	20	0.05	0.01	0	0.18
CP-NCI	lambdamart	DCG@10	20	0	0	0	0.02
EP-NCI	lambdamart	DCG@10	20	0	0	0	0.02
CP-CI	lambdamart	DCG@10	10	0.1	0.01	0	0.28
EP-CI	lambdamart	DCG@10	10	0.05	0.01	0	0.18
CP-NCI	lambdamart	DCG@10	10	0	0	0	0.02
EP-NCI	lambdamart	DCG@10	10	0.06	0	0	0.14
CP-CI	lambdamart	DCG@10	5	0.1	0.01	0	0.28

Table B.2: t_{ff} per dataset, algorithm and configuration

EP-CI	lambdamart	DCG@10	5	0.15	0.02	0	0.36
CP-NCI	lambdamart	DCG@10	5	0.04	0	0	0.14
EP-NCI	lambdamart	DCG@10	5	0.04	0	0	0.14
CP-CI	lambdamart	MAP	30	0.07	0	0.03	0.14
EP-CI	lambdamart	MAP	30	0.07	0	0.03	0.14
CP-NCI	lambdamart	MAP	30	0.07	0	0.03	0.14
EP-NCI	lambdamart	MAP	30	0.07	0	0.03	0.14
CP-CI	lambdamart	MAP	20	0.07	0	0.03	0.14
EP-CI	lambdamart	MAP	20	0.07	0	0.03	0.14
CP-NCI	lambdamart	MAP	20	0.07	0	0.03	0.14
EP-NCI	lambdamart	MAP	20	0.07	0	0.03	0.14
CP-CI	lambdamart	MAP	10	0.07	0	0.03	0.14
EP-CI	lambdamart	MAP	10	0.07	0	0.03	0.14
CP-NCI	lambdamart	MAP	10	0.07	0	0.03	0.14
EP-NCI	lambdamart	MAP	10	0.07	0	0.03	0.14
CP-CI	lambdamart	MAP	5	0.07	0	0.03	0.14
EP-CI	lambdamart	MAP	5	0.07	0	0.03	0.14
CP-NCI	lambdamart	MAP	5	0.07	0	0.03	0.14
EP-NCI	lambdamart	MAP	5	0.07	0	0.03	0.14
CP-CI	lambdamart	NDCG@20	30	0.23	0.1	0	0.93
EP-CI	lambdamart	NDCG@20	30	0.1	0.01	0.01	0.25
CP-NCI	lambdamart	NDCG@20	30	0.01	0	0	0.02
EP-NCI	lambdamart	NDCG@20	30	0.06	0	0	0.14
CP-CI	lambdamart	NDCG@20	20	0.24	0.08	0	0.79
EP-CI	lambdamart	NDCG@20	20	0.14	0.02	0	0.43
CP-NCI	lambdamart	NDCG@20	20	0	0	0	0.02
EP-NCI	lambdamart	NDCG@20	20	0.06	0	0	0.14
CP-CI	lambdamart	NDCG@20	10	0.1	0.01	0	0.28
EP-CI	lambdamart	NDCG@20	10	0.05	0	0.01	0.18
CP-NCI	lambdamart	NDCG@20	10	0.01	0	0	0.02
EP-NCI	lambdamart	NDCG@20	10	0.21	0.12	0.01	1
CP-CI	lambdamart	NDCG@20	5	0.1	0.01	0	0.28
EP-CI	lambdamart	NDCG@20	5	0.09	0	0.01	0.17
CP-NCI	lambdamart	NDCG@20	5	0.01	0	0	0.02
EP-NCI	lambdamart	NDCG@20	5	0.21	0.12	0.01	1
CP-CI	lambdamart	NDCG@30	30	0.25	0.09	0	0.89
EP-CI	lambdamart	NDCG@30	30	0.28	0.1	0.01	0.86
CP-NCI	lambdamart	NDCG@30	30	0.01	0	0	0.02
EP-NCI	lambdamart	NDCG@30	30	0.04	0	0	0.09
CP-CI	lambdamart	NDCG@30	20	0.27	0.06	0	0.7
EP-CI	lambdamart	NDCG@30	20	0.13	0.05	0.01	0.6
CP-NCI	lambdamart	NDCG@30	20	0.15	0.14	0	0.99
EP-NCI	lambdamart	NDCG@30	20	0.06	0	0	0.14
CP-CI	lambdamart	NDCG@30	10	0.1	0.01	0	0.28
EP-CI	lambdamart	NDCG@30	10	0.04	0	0	0.14
CP-NCI	lambdamart	NDCG@30	10	0.15	0.14	0	0.99
EP-NCI	lambdamart	NDCG@30	10	0.21	0.12	0.01	1
CP-CI	lambdamart	NDCG@30	5	0.1	0.01	0	0.28
EP-CI	lambdamart	NDCG@30	5	0.3	0.11	0	0.99
CP-NCI	lambdamart	NDCG@30	5	0.15	0.14	0	0.99
EP-NCI	lambdamart	NDCG@30	5	0.21	0.12	0.01	1
CP-CI	mart	NDCG@10	30	0.23	0.1	0	0.89

Table B.2: t_{ff} per dataset, algorithm and configuration

EP-CI	mart		NDCG@10	30	0.2	0.13	0.01	0.99
CP-NCI	mart		NDCG@10	30	0	0	0	0.02
EP-NCI	mart		NDCG@10	30	0.06	0	0	0.14
CP-CI	mart		NDCG@10	20	0.23	0.1	0	0.89
EP-CI	mart		NDCG@10	20	0.19	0.13	0.01	0.99
CP-NCI	mart		NDCG@10	20	0.17	0.13	0	0.99
EP-NCI	mart		NDCG@10	20	0.21	0.12	0.04	0.99
CP-CI	mart		NDCG@10	10	0.23	0.1	0	0.89
EP-CI	mart		NDCG@10	10	0.25	0.13	0.01	0.99
CP-NCI	mart		NDCG@10	10	0.21	0.12	0.04	0.99
EP-NCI	mart		NDCG@10	10	0.21	0.12	0.04	0.99
CP-CI	mart		NDCG@10	5	0.23	0.1	0	0.89
EP-CI	mart		NDCG@10	5	0.18	0.13	0.01	0.99
CP-NCI	mart		NDCG@10	5	0.21	0.12	0.04	0.99
EP-NCI	mart		NDCG@10	5	0.21	0.12	0.04	0.99
CP-CI	mart		ERR@10	30	0.23	0.1	0	0.89
EP-CI	mart		ERR@10	30	0.2	0.13	0.01	0.99
CP-NCI	mart		ERR@10	30	0	0	0	0.02
EP-NCI	mart		ERR@10	30	0.06	0	0	0.14
CP-CI	mart		ERR@10	20	0.23	0.1	0	0.89
EP-CI	mart		ERR@10	20	0.19	0.13	0.01	0.99
CP-NCI	mart		ERR@10	20	0.17	0.13	0	0.99
EP-NCI	mart		ERR@10	20	0.21	0.12	0.04	0.99
CP-CI	mart		ERR@10	10	0.23	0.1	0	0.89
EP-CI	mart		ERR@10	10	0.25	0.13	0.01	0.99
CP-NCI	mart		ERR@10	10	0.21	0.12	0.04	0.99
EP-NCI	mart		ERR@10	10	0.21	0.12	0.04	0.99
CP-CI	mart		DCG@10	30	0.23	0.1	0	0.89
EP-CI	mart		DCG@10	30	0.2	0.13	0.01	0.99
CP-NCI	mart		DCG@10	30	0	0	0	0.02
EP-NCI	mart		DCG@10	30	0.06	0	0	0.14
CP-CI	mart		DCG@10	20	0.23	0.1	0	0.89
EP-CI	mart		DCG@10	20	0.19	0.13	0.01	0.99
CP-NCI	mart		DCG@10	20	0.17	0.13	0	0.99
EP-NCI	mart		DCG@10	20	0.21	0.12	0.04	0.99
CP-CI	mart		DCG@10	10	0.23	0.1	0	0.89
EP-CI	mart		DCG@10	10	0.25	0.13	0.01	0.99
CP-NCI	mart		DCG@10	10	0.21	0.12	0.04	0.99
EP-NCI	mart		DCG@10	10	0.21	0.12	0.04	0.99
CP-CI	mart		DCG@10	5	0.23	0.1	0	0.89
EP-CI	mart		DCG@10	5	0.18	0.13	0.01	0.99
CP-NCI	mart		DCG@10	5	0.21	0.12	0.04	0.99
EP-NCI	mart		DCG@10	5	0.21	0.12	0.04	0.99
CP-CI	mart		MAP	30	0.23	0.1	0	0.89
EP-CI	mart		MAP	30	0.2	0.13	0.01	0.99
CP-NCI	mart		MAP	30	0	0	0	0.02
EP-NCI	mart		MAP	30	0.06	0	0	0.14
CP-CI	mart		MAP	20	0.23	0.1	0	0.89

Table B.2: t_{ff} per dataset, algorithm and configuration

EP-CI	mart	MAP	20	0.19	0.13	0.01	0.99
CP-NCI	mart	MAP	20	0.17	0.13	0	0.99
EP-NCI	mart	MAP	20	0.21	0.12	0.04	0.99
CP-CI	mart	MAP	10	0.23	0.1	0	0.89
EP-CI	mart	MAP	10	0.25	0.13	0.01	0.99
CP-NCI	mart	MAP	10	0.21	0.12	0.04	0.99
EP-NCI	mart	MAP	10	0.21	0.12	0.04	0.99
CP-CI	mart	MAP	5	0.23	0.1	0	0.89
EP-CI	mart	MAP	5	0.18	0.13	0.01	0.99
CP-NCI	mart	MAP	5	0.21	0.12	0.04	0.99
EP-NCI	mart	MAP	5	0.21	0.12	0.04	0.99
CP-CI	mart	NDCG@20	30	0.23	0.1	0	0.89
EP-CI	mart	NDCG@20	30	0.2	0.13	0.01	0.99
CP-NCI	mart	NDCG@20	30	0	0	0	0.02
EP-NCI	mart	NDCG@20	30	0.06	0	0	0.14
CP-CI	mart	NDCG@20	20	0.23	0.1	0	0.89
EP-CI	mart	NDCG@20	20	0.19	0.13	0.01	0.99
CP-NCI	mart	NDCG@20	20	0.17	0.13	0	0.99
EP-NCI	mart	NDCG@20	20	0.21	0.12	0.04	0.99
CP-CI	mart	NDCG@20	10	0.23	0.1	0	0.89
EP-CI	mart	NDCG@20	10	0.25	0.13	0.01	0.99
CP-NCI	mart	NDCG@20	10	0.21	0.12	0.04	0.99
EP-NCI	mart	NDCG@20	10	0.21	0.12	0.04	0.99
CP-CI	mart	NDCG@20	5	0.23	0.1	0	0.89
EP-CI	mart	NDCG@20	5	0.18	0.13	0.01	0.99
CP-NCI	mart	NDCG@20	5	0.21	0.12	0.04	0.99
EP-NCI	mart	NDCG@20	5	0.21	0.12	0.04	0.99
CP-CI	mart	NDCG@30	30	0.23	0.1	0	0.89
EP-CI	mart	NDCG@30	30	0.2	0.13	0.01	0.99
CP-NCI	mart	NDCG@30	30	0	0	0	0.02
EP-NCI	mart	NDCG@30	30	0.06	0	0	0.14
CP-CI	mart	NDCG@30	20	0.23	0.1	0	0.89
EP-CI	mart	NDCG@30	20	0.19	0.13	0.01	0.99
CP-NCI	mart	NDCG@30	20	0.17	0.13	0	0.99
EP-NCI	mart	NDCG@30	20	0.21	0.12	0.04	0.99
CP-CI	mart	NDCG@30	10	0.23	0.1	0	0.89
EP-CI	mart	NDCG@30	10	0.25	0.13	0.01	0.99
CP-NCI	mart	NDCG@30	10	0.21	0.12	0.04	0.99
EP-NCI	mart	NDCG@30	10	0.21	0.12	0.04	0.99
CP-CI	mart	NDCG@30	5	0.23	0.1	0	0.89
EP-CI	mart	NDCG@30	5	0.18	0.13	0.01	0.99
CP-NCI	mart	NDCG@30	5	0.21	0.12	0.04	0.99
EP-NCI	mart	NDCG@30	5	0.21	0.12	0.04	0.99

Table B.3: Selections induced by each configuration

Dataset	Algorithm	Training metric	# of trees	50-SEL		80-SEL		S-SEL	
				Size	Execution Time	Size	Execution Time	Size	Execution Time
EP-NCI	coordinateascent	NDCG@10		10	0.03	40	0.1	40	0.1
CP-CI	coordinateascent	NDCG@10		10	0.16	80	0.81	90	0.92
CP-NCI	coordinateascent	NDCG@10		10	0.03	40	0.1	40	0.1
EP-CI	coordinateascent	NDCG@10		30	0.22	90	0.77	90	0.77
EP-NCI	coordinateascent	DCG@10		10	0.03	40	0.1	40	0.1
CP-CI	coordinateascent	DCG@10		10	0.17	80	0.82	90	0.9
CP-NCI	coordinateascent	DCG@10		10	0.03	40	0.1	40	0.1
EP-CI	coordinateascent	DCG@10		30	0.22	80	0.84	100	1
EP-NCI	coordinateascent	MAP		10	0.03	40	0.1	40	0.1
CP-CI	coordinateascent	MAP		60	0.63	90	0.9	90	0.9
CP-NCI	coordinateascent	MAP		10	0.03	40	0.1	40	0.1
EP-CI	coordinateascent	MAP		60	0.63	90	0.9	90	0.9
EP-NCI	coordinateascent	NDCG@20		10	0.03	40	0.1	40	0.1
CP-CI	coordinateascent	NDCG@20		10	0.15	80	0.83	90	0.9
CP-NCI	coordinateascent	NDCG@20		10	0.03	40	0.1	40	0.1
EP-CI	coordinateascent	NDCG@20		30	0.22	90	0.77	90	0.77
EP-NCI	coordinateascent	NDCG@30		10	0.03	40	0.1	40	0.1
CP-CI	coordinateascent	NDCG@30		10	0.17	80	0.82	90	0.9
CP-NCI	coordinateascent	NDCG@30		10	0.03	40	0.1	40	0.1
EP-CI	coordinateascent	NDCG@30		10	0.09	30	0.2	50	0.4
EP-NCI	lambdamart	NDCG@10	30	40	0.09	100	1	100	1
CP-CI	lambdamart	NDCG@10	30	50	0.56	90	0.91	90	0.91
CP-NCI	lambdamart	NDCG@10	30	10	0.03	40	0.1	40	0.1
EP-CI	lambdamart	NDCG@10	30	40	0.09	40	0.09	40	0.09
EP-NCI	lambdamart	NDCG@10	20	40	0.09	100	1	100	1
CP-CI	lambdamart	NDCG@10	20	50	0.56	50	0.56	60	0.6
CP-NCI	lambdamart	NDCG@10	20	10	0.03	40	0.1	40	0.1
EP-CI	lambdamart	NDCG@10	20	30	0.17	60	0.54	70	0.73
EP-NCI	lambdamart	NDCG@10	10	30	0.09	40	0.1	100	1
CP-CI	lambdamart	NDCG@10	10	40	0.15	40	0.15	40	0.15
CP-NCI	lambdamart	NDCG@10	10	10	0.03	40	0.1	40	0.1
EP-CI	lambdamart	NDCG@10	10	40	0.23	60	0.53	70	0.72
EP-NCI	lambdamart	NDCG@10	5	40	0.09	100	1	100	1
CP-CI	lambdamart	NDCG@10	5	40	0.15	40	0.15	40	0.15
CP-NCI	lambdamart	NDCG@10	5	10	0.03	40	0.1	40	0.1
EP-CI	lambdamart	NDCG@10	5	40	0.3	50	0.41	70	0.74
EP-NCI	lambdamart	DCG@10	30	40	0.09	100	1	100	1
CP-CI	lambdamart	DCG@10	30	30	0.41	50	0.56	60	0.61
CP-NCI	lambdamart	DCG@10	30	10	0.03	40	0.1	40	0.1
EP-CI	lambdamart	DCG@10	30	40	0.2	70	0.56	80	0.74
EP-NCI	lambdamart	DCG@10	20	10	0.03	40	0.1	40	0.1
CP-CI	lambdamart	DCG@10	20	40	0.15	40	0.15	40	0.15
CP-NCI	lambdamart	DCG@10	20	10	0.03	40	0.1	40	0.1
EP-CI	lambdamart	DCG@10	20	40	0.2	70	0.56	80	0.74
EP-NCI	lambdamart	DCG@10	10	40	0.09	100	1	100	1
CP-CI	lambdamart	DCG@10	10	40	0.15	40	0.15	40	0.15
CP-NCI	lambdamart	DCG@10	10	10	0.03	40	0.1	40	0.1
EP-CI	lambdamart	DCG@10	10	40	0.2	70	0.56	80	0.74
EP-NCI	lambdamart	DCG@10	5	40	0.09	100	1	100	1

Table B.3: Selections induced by each configuration

CP-CI	lambdamart	DCG@10	5	40	0.15	40	0.15	40	0.15
CP-NCI	lambdamart	DCG@10	5	40	0.09	40	0.09	40	0.09
EP-CI	lambdamart	DCG@10	5	50	0.21	70	0.57	90	0.86
EP-NCI	lambdamart	MAP	30	40	0.09	40	0.09	40	0.09
CP-CI	lambdamart	MAP	30	40	0.09	40	0.09	40	0.09
CP-NCI	lambdamart	MAP	30	40	0.09	40	0.09	40	0.09
EP-CI	lambdamart	MAP	30	40	0.09	40	0.09	40	0.09
EP-NCI	lambdamart	MAP	20	40	0.09	40	0.09	40	0.09
CP-CI	lambdamart	MAP	20	40	0.09	40	0.09	40	0.09
CP-NCI	lambdamart	MAP	20	40	0.09	40	0.09	40	0.09
EP-CI	lambdamart	MAP	20	40	0.09	40	0.09	40	0.09
EP-NCI	lambdamart	MAP	10	40	0.09	40	0.09	40	0.09
CP-CI	lambdamart	MAP	10	40	0.09	40	0.09	40	0.09
CP-NCI	lambdamart	MAP	10	40	0.09	40	0.09	40	0.09
EP-CI	lambdamart	MAP	10	40	0.09	40	0.09	40	0.09
EP-NCI	lambdamart	MAP	5	40	0.09	40	0.09	40	0.09
CP-CI	lambdamart	MAP	5	40	0.09	40	0.09	40	0.09
CP-NCI	lambdamart	MAP	5	40	0.09	40	0.09	40	0.09
EP-CI	lambdamart	MAP	5	40	0.09	40	0.09	40	0.09
EP-NCI	lambdamart	NDCG@20	30	40	0.09	100	1	100	1
CP-CI	lambdamart	NDCG@20	30	30	0.44	70	0.73	90	0.92
CP-NCI	lambdamart	NDCG@20	30	10	0.03	40	0.1	40	0.1
EP-CI	lambdamart	NDCG@20	30	40	0.22	60	0.45	70	0.67
EP-NCI	lambdamart	NDCG@20	20	40	0.09	100	1	100	1
CP-CI	lambdamart	NDCG@20	20	50	0.59	70	0.7	70	0.7
CP-NCI	lambdamart	NDCG@20	20	10	0.03	40	0.1	40	0.1
EP-CI	lambdamart	NDCG@20	20	60	0.47	60	0.47	70	0.66
EP-NCI	lambdamart	NDCG@20	10	40	0.09	100	1	100	1
CP-CI	lambdamart	NDCG@20	10	40	0.15	40	0.15	40	0.15
CP-NCI	lambdamart	NDCG@20	10	10	0.03	40	0.1	40	0.1
EP-CI	lambdamart	NDCG@20	10	40	0.22	60	0.51	80	0.83
EP-NCI	lambdamart	NDCG@20	5	40	0.09	100	1	100	1
CP-CI	lambdamart	NDCG@20	5	40	0.15	40	0.15	40	0.15
CP-NCI	lambdamart	NDCG@20	5	10	0.03	40	0.1	40	0.1
EP-CI	lambdamart	NDCG@20	5	40	0.24	60	0.53	70	0.71
EP-NCI	lambdamart	NDCG@30	30	40	0.09	100	1	100	1
CP-CI	lambdamart	NDCG@30	30	30	0.41	70	0.73	90	0.91
CP-NCI	lambdamart	NDCG@30	30	10	0.03	40	0.1	40	0.1
EP-CI	lambdamart	NDCG@30	30	50	0.36	70	0.67	90	0.92
EP-NCI	lambdamart	NDCG@30	20	40	0.09	100	1	100	1
CP-CI	lambdamart	NDCG@30	20	50	0.56	50	0.56	60	0.6
CP-NCI	lambdamart	NDCG@30	20	40	0.1	40	0.1	100	1
EP-CI	lambdamart	NDCG@30	20	50	0.36	70	0.71	70	0.71
EP-NCI	lambdamart	NDCG@30	10	40	0.09	100	1	100	1
CP-CI	lambdamart	NDCG@30	10	40	0.15	40	0.15	40	0.15
CP-NCI	lambdamart	NDCG@30	10	40	0.1	40	0.1	100	1
EP-CI	lambdamart	NDCG@30	10	40	0.09	40	0.09	40	0.09
EP-NCI	lambdamart	NDCG@30	5	40	0.09	100	1	100	1
CP-CI	lambdamart	NDCG@30	5	40	0.15	40	0.15	40	0.15
CP-NCI	lambdamart	NDCG@30	5	40	0.1	40	0.1	100	1
EP-CI	lambdamart	NDCG@30	5	50	0.42	50	0.42	100	1
EP-NCI	mart	NDCG@10	30	40	0.09	40	0.09	100	1

Table B.3: Selections induced by each configuration

CP-CI	mart	NDCG@10	30	20	0.18	70	0.84	100	1
CP-NCI	mart	NDCG@10	30	10	0.03	40	0.1	40	0.1
EP-CI	mart	NDCG@10	30	70	0.55	90	0.86	100	1
EP-NCI	mart	NDCG@10	20	40	0.09	40	0.09	100	1
CP-CI	mart	NDCG@10	20	20	0.18	70	0.84	100	1
CP-NCI	mart	NDCG@10	20	40	0.1	40	0.1	100	1
EP-CI	mart	NDCG@10	20	60	0.43	80	0.8	100	1
EP-NCI	mart	NDCG@10	10	40	0.09	40	0.09	100	1
CP-CI	mart	NDCG@10	10	20	0.17	70	0.76	90	0.92
CP-NCI	mart	NDCG@10	10	40	0.09	40	0.09	100	1
EP-CI	mart	NDCG@10	10	60	0.41	60	0.41	100	1
EP-NCI	mart	NDCG@10	5	40	0.09	40	0.09	100	1
CP-CI	mart	NDCG@10	5	20	0.17	70	0.76	90	0.92
CP-NCI	mart	NDCG@10	5	40	0.09	40	0.09	100	1
EP-CI	mart	NDCG@10	5	50	0.25	60	0.4	100	1
EP-NCI	mart	ERR@10	30	40	0.09	40	0.09	100	1
CP-CI	mart	ERR@10	30	20	0.18	70	0.84	100	1
CP-NCI	mart	ERR@10	30	10	0.03	40	0.1	40	0.1
EP-CI	mart	ERR@10	30	70	0.55	90	0.86	100	1
EP-NCI	mart	ERR@10	20	40	0.09	40	0.09	100	1
CP-CI	mart	ERR@10	20	20	0.18	70	0.84	100	1
CP-NCI	mart	ERR@10	20	40	0.1	40	0.1	100	1
EP-CI	mart	ERR@10	20	60	0.43	80	0.8	100	1
EP-NCI	mart	ERR@10	10	40	0.09	40	0.09	100	1
CP-CI	mart	ERR@10	10	20	0.17	70	0.76	90	0.92
CP-NCI	mart	ERR@10	10	40	0.09	40	0.09	100	1
EP-CI	mart	ERR@10	10	60	0.41	60	0.41	100	1
EP-NCI	mart	ERR@10	5	40	0.09	40	0.09	100	1
CP-CI	mart	ERR@10	5	20	0.17	70	0.76	90	0.92
CP-NCI	mart	ERR@10	5	40	0.09	40	0.09	100	1
EP-CI	mart	ERR@10	5	50	0.25	60	0.4	100	1
EP-NCI	mart	DCG@10	30	40	0.09	40	0.09	100	1
CP-CI	mart	DCG@10	30	20	0.18	70	0.84	100	1
CP-NCI	mart	DCG@10	30	10	0.03	40	0.1	40	0.1
EP-CI	mart	DCG@10	30	70	0.55	90	0.86	100	1
EP-NCI	mart	DCG@10	20	40	0.09	40	0.09	100	1
CP-CI	mart	DCG@10	20	20	0.18	70	0.84	100	1
CP-NCI	mart	DCG@10	20	40	0.1	40	0.1	100	1
EP-CI	mart	DCG@10	20	60	0.43	80	0.8	100	1
EP-NCI	mart	DCG@10	10	40	0.09	40	0.09	100	1
CP-CI	mart	DCG@10	10	20	0.17	70	0.76	90	0.92
CP-NCI	mart	DCG@10	10	40	0.09	40	0.09	100	1
EP-CI	mart	DCG@10	10	60	0.41	60	0.41	100	1
EP-NCI	mart	DCG@10	5	40	0.09	40	0.09	100	1
CP-CI	mart	DCG@10	5	20	0.17	70	0.76	90	0.92
CP-NCI	mart	DCG@10	5	40	0.09	40	0.09	100	1
EP-CI	mart	DCG@10	5	50	0.25	60	0.4	100	1
EP-NCI	mart	MAP	30	40	0.09	40	0.09	100	1
CP-CI	mart	MAP	30	20	0.18	70	0.84	100	1
CP-NCI	mart	MAP	30	10	0.03	40	0.1	40	0.1
EP-CI	mart	MAP	30	70	0.55	90	0.86	100	1
EP-NCI	mart	MAP	20	40	0.09	40	0.09	100	1

Table B.3: Selections induced by each configuration

CP-CI	mart	MAP	20	20	0.18	70	0.84	100	1
CP-NCI	mart	MAP	20	40	0.1	40	0.1	100	1
EP-CI	mart	MAP	20	60	0.43	80	0.8	100	1
EP-NCI	mart	MAP	10	40	0.09	40	0.09	100	1
CP-CI	mart	MAP	10	20	0.17	70	0.76	90	0.92
CP-NCI	mart	MAP	10	40	0.09	40	0.09	100	1
EP-CI	mart	MAP	10	60	0.41	60	0.41	100	1
EP-NCI	mart	MAP	5	40	0.09	40	0.09	100	1
CP-CI	mart	MAP	5	20	0.17	70	0.76	90	0.92
CP-NCI	mart	MAP	5	40	0.09	40	0.09	100	1
EP-CI	mart	MAP	5	50	0.25	60	0.4	100	1
EP-NCI	mart	NDCG@20	30	40	0.09	40	0.09	100	1
CP-CI	mart	NDCG@20	30	20	0.18	70	0.84	100	1
CP-NCI	mart	NDCG@20	30	10	0.03	40	0.1	40	0.1
EP-CI	mart	NDCG@20	30	70	0.55	90	0.86	100	1
EP-NCI	mart	NDCG@20	20	40	0.09	40	0.09	100	1
CP-CI	mart	NDCG@20	20	20	0.18	70	0.84	100	1
CP-NCI	mart	NDCG@20	20	40	0.1	40	0.1	100	1
EP-CI	mart	NDCG@20	20	60	0.43	80	0.8	100	1
EP-NCI	mart	NDCG@20	10	40	0.09	40	0.09	100	1
CP-CI	mart	NDCG@20	10	20	0.17	70	0.76	90	0.92
CP-NCI	mart	NDCG@20	10	40	0.09	40	0.09	100	1
EP-CI	mart	NDCG@20	10	60	0.41	60	0.41	100	1
EP-NCI	mart	NDCG@20	5	40	0.09	40	0.09	100	1
CP-CI	mart	NDCG@20	5	20	0.17	70	0.76	90	0.92
CP-NCI	mart	NDCG@20	5	40	0.09	40	0.09	100	1
EP-CI	mart	NDCG@20	5	50	0.25	60	0.4	100	1
EP-NCI	mart	NDCG@30	30	40	0.09	40	0.09	100	1
CP-CI	mart	NDCG@30	30	20	0.18	70	0.84	100	1
CP-NCI	mart	NDCG@30	30	10	0.03	40	0.1	40	0.1
EP-CI	mart	NDCG@30	30	70	0.55	90	0.86	100	1
EP-NCI	mart	NDCG@30	20	40	0.09	40	0.09	100	1
CP-CI	mart	NDCG@30	20	20	0.18	70	0.84	100	1
CP-NCI	mart	NDCG@30	20	40	0.1	40	0.1	100	1
EP-CI	mart	NDCG@30	20	60	0.43	80	0.8	100	1
EP-NCI	mart	NDCG@30	10	40	0.09	40	0.09	100	1
CP-CI	mart	NDCG@30	10	20	0.17	70	0.76	90	0.92
CP-NCI	mart	NDCG@30	10	40	0.09	40	0.09	100	1
EP-CI	mart	NDCG@30	10	60	0.41	60	0.41	100	1
EP-NCI	mart	NDCG@30	5	40	0.09	40	0.09	100	1
CP-CI	mart	NDCG@30	5	20	0.17	70	0.76	90	0.92
CP-NCI	mart	NDCG@30	5	40	0.09	40	0.09	100	1
EP-CI	mart	NDCG@30	5	50	0.25	60	0.4	100	1